



WELCOME



MongoDB Atlas

Developer Data Platform Workshop



David Hiltenbrand - Solutions Architect
Memphis, TN



Soheyl Rafi - Solutions Architect
Atlanta, GA



Ian Paeth - Enterprise Account Executive
Atlanta, GA

Agenda

Setting up Atlas

Downloading Compass

CRUD Operations

Aggregations

Charts/ Search / Vector Search



Repo

<https://github.com/srafimdb/AtlasWorkshop/blob/main/README.md>

or

<https://shorturl.at/lxFU1>



Now it's your turn to try
MongoDB Atlas



Exercise One

MongoDB Atlas Demo

DEPLOYMENT

Database

Data Lake PREVIEW

DATA SERVICES

Triggers

Data API

Data Federation

SECURITY

Database Access

Network Access

Advanced

New On Atlas 4

RPC FIELD M

Data

Database

User Name

DO_NOC

Add New Database User

Create a database user to grant an application or user, access to databases and collections in your clusters in this Atlas project. Granular access control can be configured with default privileges or custom roles. You can grant access to an Atlas project or organization using the corresponding [Access Manager](#)

Authentication Method

Password

Certificate

AWS IAM
(MongoDB 4.4 and up)

MongoDB uses [SCRAM](#) as its default authentication method.

Password Authentication

appUser

.....

SHOW

[Autogenerate Secure Password](#)

[Copy](#)

Database User Privileges

Configure role based access control by assigning database users a mix of one built-in role, multiple custom roles, and multiple specific privileges. A user will gain access to all actions within the roles assigned to them, not just the actions those roles share in common. **You must choose at least one role or privilege.** [Learn more about roles.](#)

Built-in Role

Select one [built-in role](#) for this user.

Read and write to any database

1 SELECTED



Your changes (current action):

Access

Peering

Private

Add IP Access List Entry

Atlas only allows client connections to a cluster from entries in the project's IP Access List. Each entry should either be a single IP address, a CIDR-notated range of addresses, or an AWS Security Group ID. [Learn more](#).

[ADD CURRENT IP ADDRESS](#)[ALLOW ACCESS FROM ANYWHERE](#)**Access List Entry:**

93.231.182.61

Comment:

Optional comment describing this entry



This entry is temporary and will be deleted in

6 hours ▾

[Cancel](#)[Confirm](#)

Add an IP address

Configure which IP addresses can access your cluster.

[Add IP Address](#)[Learn more](#)



Load Sample Dataset

The screenshot shows the MongoDB Atlas Cluster Overview page for Cluster0. The top navigation bar includes 'Cluster0', 'Connect', 'View Monitoring', 'Browse Collections', and a '...' button. A modal window titled 'Enhance Your Experience' offers an upgrade to a M30 cluster. The main cluster details table shows the following information:

VERSION	REGION	CLUSTER TIER	TYPE	BACKUPS	LINKED APP SERVICES	ATLAS SQL	ONLINE ARCHIVE	ATLAS SEARCH
7.0	AWS / N. Virginia (us-east-1)	M10 (General)	Replica Set - 3 nodes	Active	None Linked	Connect	Active	8 search indexes

Below the table is a '+ Add Tag' button. On the right side of the screen, there are monitoring charts for network traffic and disk usage, and a 'DEDICATED' status indicator. A context menu is open over the cluster name, listing options: 'Edit Configuration', 'Command Line Tools', 'Load Sample Dataset' (which is highlighted with a red box), 'Download Logs', 'Test Resilience', 'Take Snapshot Now', 'Pause Cluster', and 'Terminate'.



Download Compass



<https://www.mongodb.com/try/download/compass>

MongoDB Compass

Easily explore and manipulate your database with Compass, the GUI for MongoDB. Intuitive and flexible, Compass provides detailed schema visualizations, real-time performance metrics, sophisticated querying abilities, and much more.

Please note that MongoDB Compass comes in three versions: **a full version** with all features, **a read-only version** without write or delete capabilities, and **an isolated edition**, whose sole network connection is to the MongoDB instance.

For more information, see our [documentation pages](#).

○ Compass
The full version of MongoDB Compass, with all features and capabilities.

○ Readonly Edition
This version is limited strictly to read operations, with all write and delete capabilities removed.

○ Isolated Edition
This version disables all network connections except the connection to the MongoDB instance.

Available Downloads

Version 1.33.1 (Stable) ▼

Platform macOS 64-bit (10.14+) ▼

Package dmg

Download Copy Link

[Documentation](#) [Archived releases](#)

ACTIVE

Atlas

Connect to Cluster1

✓ Setup connection security ✔ Choose a connection method ✖ Connect ✖

Choose a connection method [View documentation](#)

Get your pre-formatted connection string by selecting your tool below.



Connect with the MongoDB Shell

Interact with your cluster using MongoDB's interactive Javascript interface



Connect your application

Connect your application to your cluster using MongoDB's native drivers



Connect using MongoDB Compass

Explore, modify, and visualize your data with MongoDB's GUI



Connect using VS Code

Connect to a MongoDB host in Visual Studio Code



Connect Your Business Intelligence Tool ([Enable the BI Connector](#))

Get connection settings for Business Intelligence tools and visualize your data

[Go Back](#)[Close](#)



Exercise Two

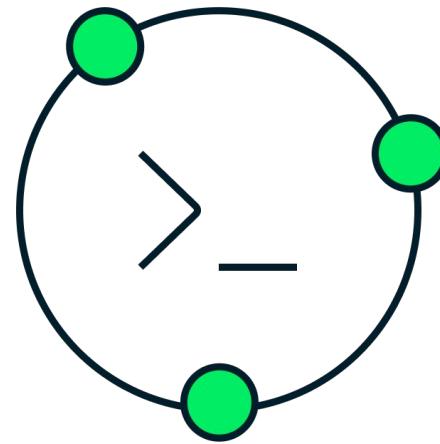
MongoDB Data GUI & CRUD



Developer tools ecosystem



Build Applications
IDE Integrations
VS Code for MongoDB



Navigate & explore your data
MongoDB Compass



Interface with MongoDB
MongoDB Shell



MongoDB Query API: rich and expressive

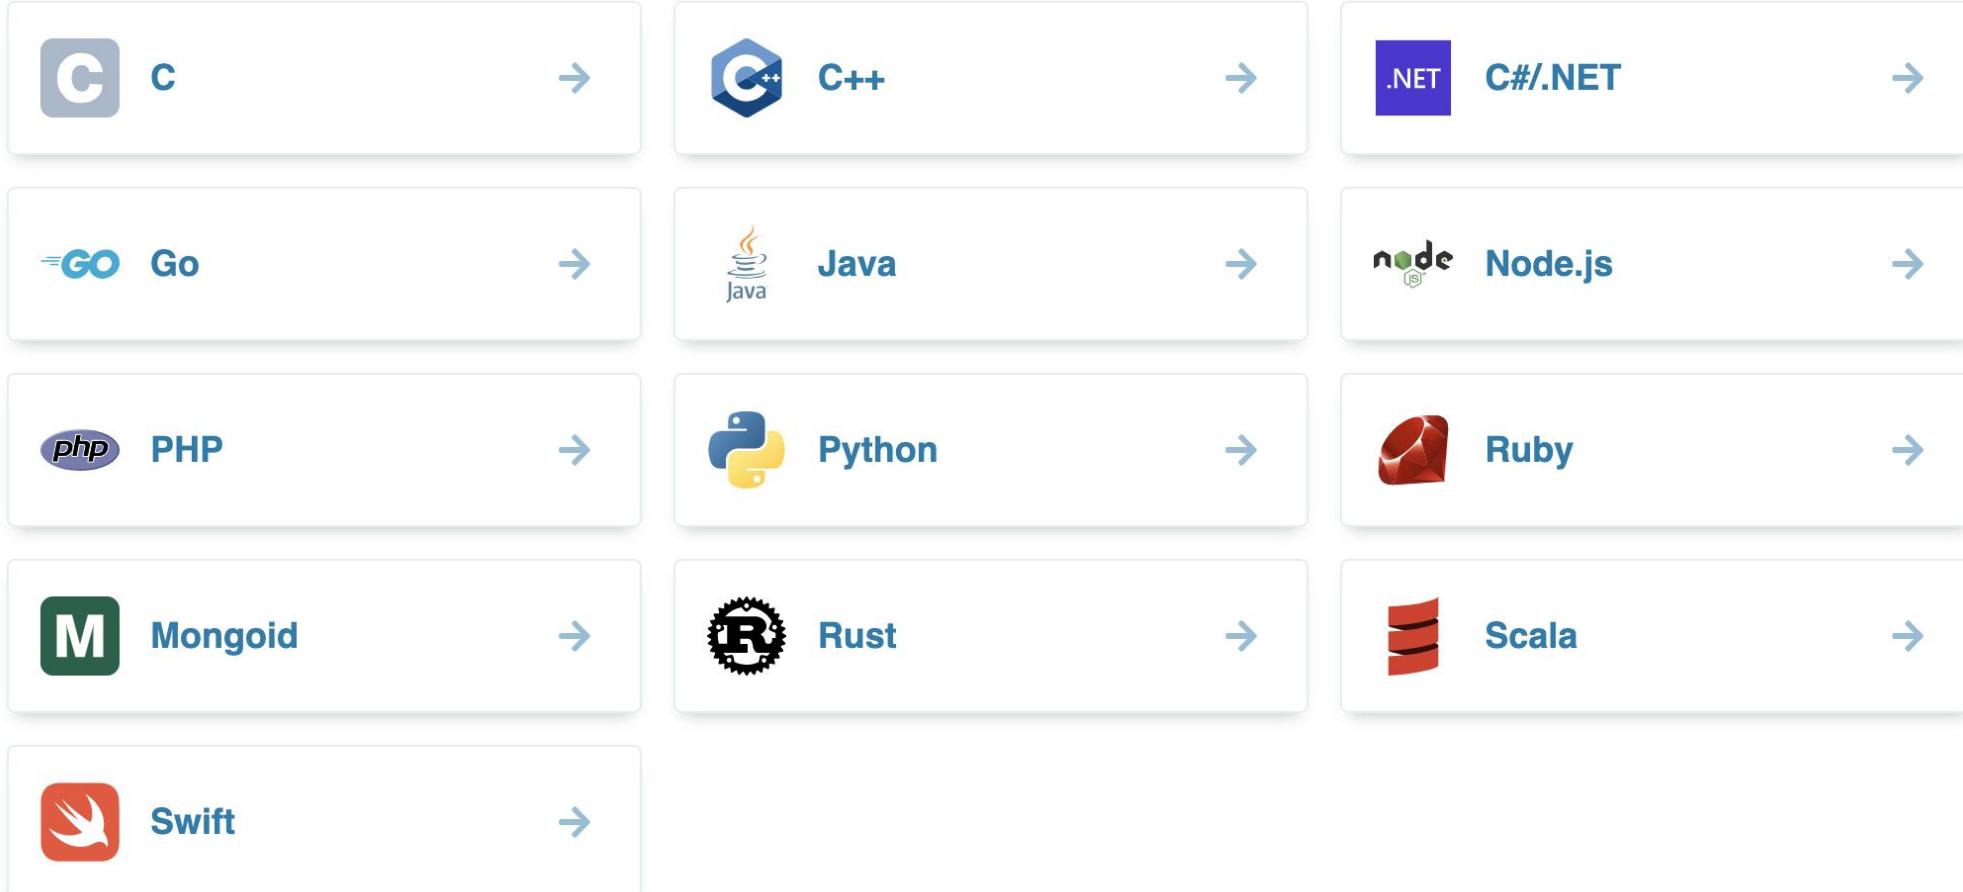
Expressive queries	<ul style="list-style-type: none">Find anyone with phone # “1-212...”Check if the person with number “555...” is on the “do not call” list
Geospatial	<ul style="list-style-type: none">Find the best offer for the customer at geo coordinates of 42nd St. and 6th Ave
Text search	<ul style="list-style-type: none">Find all tweets that mention the firm within the last 2 days
Aggregation	<ul style="list-style-type: none">Count and sort number of customers by city, compute min, max, and average spend
Native binary JSON support	<ul style="list-style-type: none">Add an additional phone number to Mark Smith’s record without rewriting the documentUpdate just 2 phone numbers out of 10Sort on the modified date
JOIN (\$lookup)	<ul style="list-style-type: none">Query for all San Francisco residences, lookup their transactions, and sum the amount by person
Graph queries (\$graphLookup)	<ul style="list-style-type: none">Query for all people within 3 degrees of separation from Mark

MongoDB

```
{  
  customer_id : 1,  
  first_name : "Mark",  
  last_name : "Smith",  
  city : "San Francisco",  
  phones: [  
    {  
      number : "1-212-777-1212",  
      type : "work"  
    },  
    {  
      number : "1-212-777-1213",  
      type : "cell"  
    }  
  ]  
}
```



Intuitive: client drivers

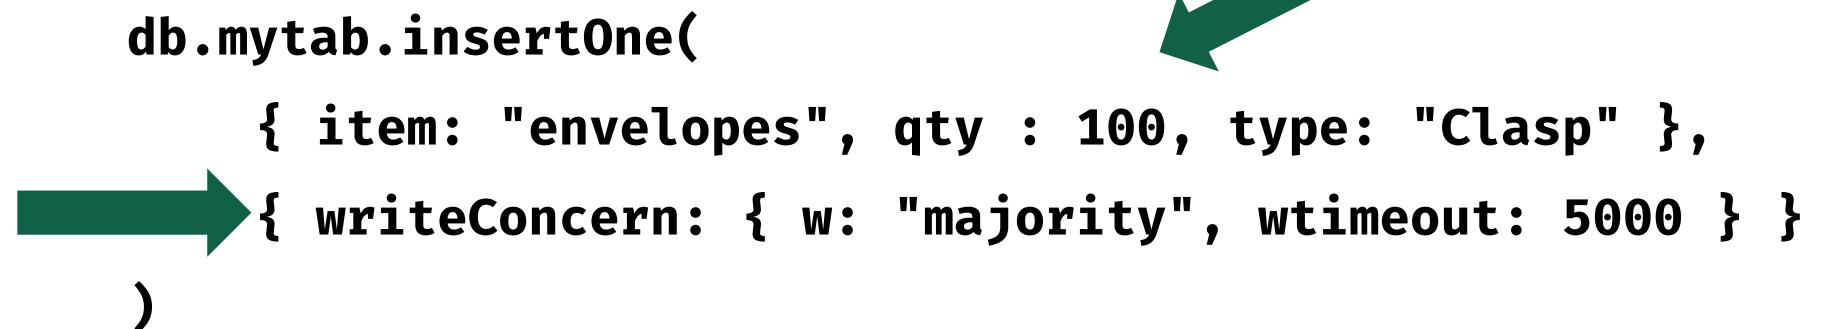


- Common CRUD capabilities but idiomatic to each language
- Uniform HA & Failover capabilities across all

insertOne()

- Inserts document(s) into a collection
- Can optionally define the writeConcern of the operation
- `_id` field is automatically generated with an ObjectId

```
db.mytab.insertOne(  
  { item: "envelopes", qty : 100, type: "Clasp" },  
  { writeConcern: { w: "majority", wtimeout: 5000 } }  
)
```



Optional
writeConcern
Argument

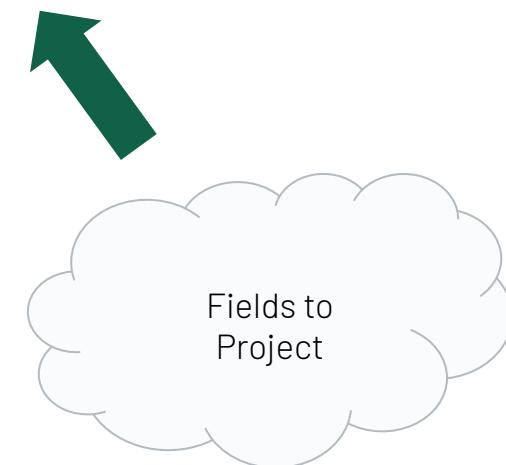
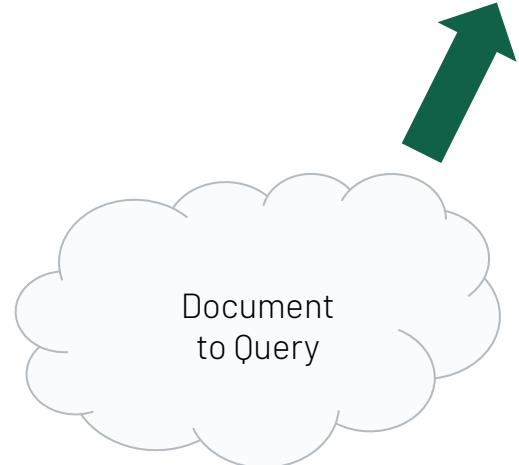
Document
to Insert



`find()` and `findOne()`

- Get documents from the collection (table)
- Can filter by content (query) and by what is returned (project)
- `findOne()` returns just the first document

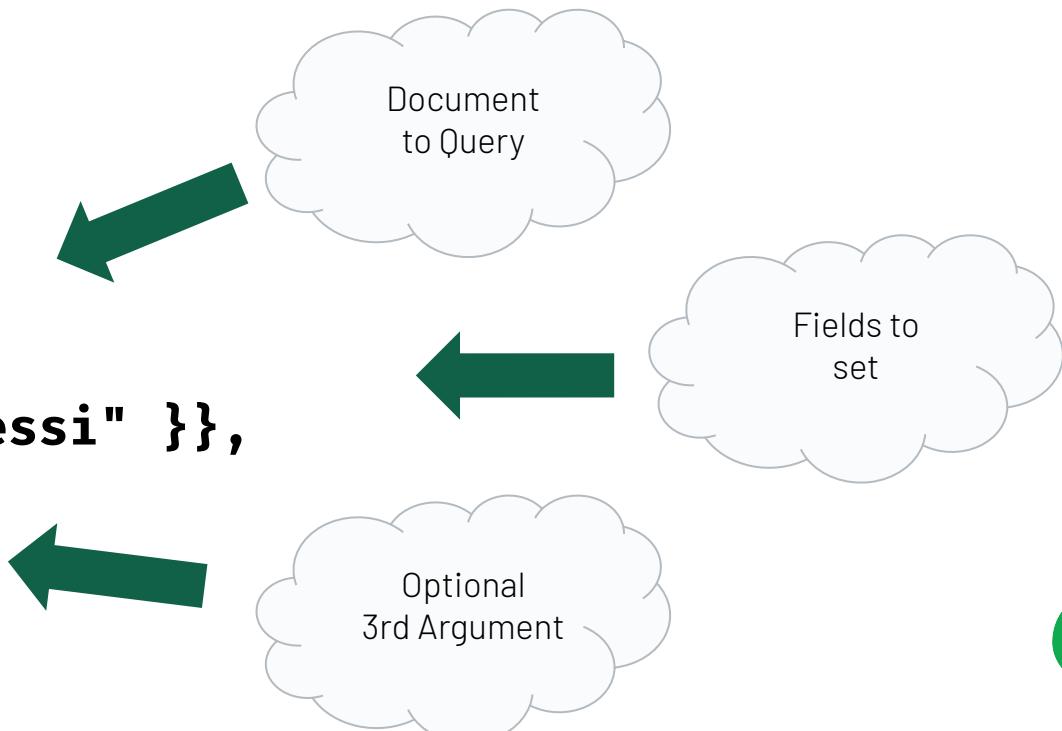
```
db.mytab.find({x:1, y:2}, {a:1, b:1, c:1})
```



updateMany() and updateOne()

- Modifies existing document(s) in a collection
- Use with a field update operator (MQL)
- Query for the documents you want to update
- Can optionally upsert, multi-update, define writeConcern, etc.

```
db.books.updateOne(  
  { author: "Roger Federer" },  
  { $set: { author: "Lionel Messi" }},  
  { multi: true }  
)
```



deleteMany()

- Deletes a document from the collection
- Takes an optional writeConcern argument

```
db.products.deleteMany( { qty: { $gt: 20 } },  
{ writeConcern: { w: "majority", wtimeout: 5000 } })
```



Comparison Query Operators

- \$lt : Exists and less than
- \$lte : Exists and less than or equal to
- \$gt : Exists and greater than
- \$gte : Exists and greater than or equal to
- \$ne : Does not exist or does but not equal to
- \$in : Exists and in a set
- \$nin : Does not exist or not in a set

```
{ age : { $gte : 21 }}
```

```
{ temp : { $gt: 10, $lt: 30 }}
```



Field Update Operators (well some...)

- \$currentDate : Sets the value of a field to current date
- \$inc : Increments value
- \$min : Update if field is less than \$min value
- \$max : Update if field is greater than \$max value
- \$mul : Multiply field value
- \$set : Set value of field
- \$unset : Removes field from document

```
{ $set : { bestDB : "MongoDB" } }  
{ $mul: { amount : 5, ... } }
```





We will now be using
sample_airbnb.listingsAndReviews
& sample_mflix.movies
database/collection from now on



JEOPARDY!

Please answer via the **Chat**



We will do the first one together now.





CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987
db.movies.find({genres: "Comedy"})	



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987
db.movies.find({genres: "Comedy"})	“Comedy” as one of their genres



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987
db.movies.find({genres: "Comedy"})	“Comedy” as one of their genres
db.movies.find({genres: ["Comedy"]})	



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987
db.movies.find({genres: "Comedy"})	“Comedy” as one of their genres
db.movies.find({genres: ["Comedy"]})	“Comedy” as the <i>only</i> genre



CRUD: Jeopardy

Query	What's the Question?
<code>db.movies.find({year: 1987})</code>	... from 1987
<code>db.movies.find({genres: "Comedy"})</code>	“Comedy” as one of their genres
<code>db.movies.find({genres: ["Comedy"]})</code>	“Comedy” as the <i>only</i> genre
<code>db.movies.find({genres: {\$in:["Comedy", "Drama"]}})</code>	



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987
db.movies.find({genres: "Comedy"})	“Comedy” as one of their genres
db.movies.find({genres: ["Comedy"]})	“Comedy” as the <i>only</i> genre
db.movies.find({genres: {\$in:["Comedy", "Drama"]}})	“Comedy” or “Drama” as the genre



CRUD: Jeopardy

Query	What's the Question?
<code>db.movies.find({year: 1987})</code>	... from 1987
<code>db.movies.find({genres: "Comedy"})</code>	“Comedy” as one of their genres
<code>db.movies.find({genres: ["Comedy"]})</code>	“Comedy” as the <i>only</i> genre
<code>db.movies.find({genres: {\$in:["Comedy", "Drama"]}})</code>	“Comedy” or “Drama” as the genre
<code>db.movies.find({revenues: {\$exists: true}})</code>	



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987
db.movies.find({genres: "Comedy"})	“Comedy” as one of their genres
db.movies.find({genres: ["Comedy"]})	“Comedy” as the <i>only</i> genre
db.movies.find({genres: {\$in:["Comedy", "Drama"]}})	“Comedy” or “Drama” as the genre
db.movies.find({revenues: {\$exists: true}})	Documents that have a revenue field



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987
db.movies.find({genres: "Comedy"})	“Comedy” as one of their genres
db.movies.find({genres: ["Comedy"]})	“Comedy” as the <i>only</i> genre
db.movies.find({genres: {\$in:["Comedy", "Drama"]}})	“Comedy” or “Drama” as the genre
db.movies.find({revenues: {\$exists: true}})	Documents that have a revenue field
db.movies.find({imdb.rating: {\$gt: 8.0}, rated: "PG"})	



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987
db.movies.find({genres: "Comedy"})	“Comedy” as one of their genres
db.movies.find({genres: ["Comedy"]})	“Comedy” as the <i>only</i> genre
db.movies.find({genres: {\$in:["Comedy", "Drama"]}})	“Comedy” or “Drama” as the genre
db.movies.find({revenues: {\$exists: true}})	Documents that have a revenue field
db.movies.find({imdb.rating: {\$gt: 8.0}, rated: "PG"})	IMDB Rating > 8.0 and PG Rating



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987
db.movies.find({genres: "Comedy"})	“Comedy” as one of their genres
db.movies.find({genres: ["Comedy"]})	“Comedy” as the <i>only</i> genre
db.movies.find({genres: {\$in:["Comedy", "Drama"]}})	“Comedy” or “Drama” as the genre
db.movies.find({revenues: {\$exists: true}})	Documents that have a revenue field
db.movies.find({imdb.rating: {\$gt: 8.0}, rated: "PG"})	IMDB Rating > 8.0 and PG Rating
db.movies.find({title: {\$regex: '^Dr. Strangelove'}})	



CRUD: Jeopardy

Query	What's the Question?
db.movies.find({year: 1987})	... from 1987
db.movies.find({genres: "Comedy"})	“Comedy” as one of their genres
db.movies.find({genres: ["Comedy"]})	“Comedy” as the <i>only</i> genre
db.movies.find({genres: {\$in:["Comedy", "Drama"]}})	“Comedy” or “Drama” as the genre
db.movies.find({revenues: {\$exists: true}})	Documents that have a revenue field
db.movies.find({imdb.rating: {\$gt: 8.0}, rated: "PG"})	IMDB Rating > 8.0 and PG Rating
db.movies.find({title: {\$regex: '^Dr. Strangelove'}})	Title starting with “Dr. Strangelove”



<https://shorturl.at/lxFU1>

Hands-On: Your turn

sample_airbnb.listingsAndReviews

5.6k
DOCUMENTS

4
INDEXES

Documents

Aggregations

Schema

Explain Plan

Indexes

Validation

Filter



{"address.market": "Hong Kong", "property_type" : "Apartment"}

Reset

Find



More Options

ADD DATA

EXPORT COLLECTION

1 - 20 of 444



```
_id: "10893326"
listing_url: "https://www.airbnb.com/rooms/10893326"
name: "Designer Apartment in Sheung Wan"
summary: "We live in a beautiful and spacious one bedroom apartment in the heart..."
space: ""
description: "We live in a beautiful and spacious one bedroom apartment in the heart..."
neighborhood_overview: "Often described as the new SOHO, there are chic coffee shops and inter..."
notes: ""
transit: "The area is well served by minibuses, City Bus, taxis, and the old sty..."
access: ""
interaction: ""
house_rules: ""
property_type: "Apartment"
room_type: "Entire home/apt"
bed_type: "Real Bed"
minimum_nights: "1"
maximum_nights: "1125"
cancellation_policy: "strict_14_with_grace_period"
last_scraped: 2019-03-11T04:00:00.000+00:00
calendar_last_scraped: 2019-03-11T04:00:00.000+00:00
first_review: 2016-03-06T05:00:00.000+00:00
last_review: 2019-02-11T05:00:00.000+00:00
accommodates: 2
```

```
use sample_airbnb
db.listingsAndReviews.find({ "address.market": "Hong Kong", "property_type" : "Apartment" })
```

sample_airbnb.listingsAndReviews

5.6k

DOCUMENTS

4

INDEXES

[Documents](#)[Aggregations](#)[Schema](#)[Explain Plan](#)[Indexes](#)[Validation](#)

Filter



{ "bedrooms": {"\$gte": 3}, "price": {"\$gte": 1400, "\$lte": 1500}}

Reset

Find



More Options ▾

[ADD DATA](#)[EXPORT COLLECTION](#)

1 - 7 of 7



```
_id: "1176693"
listing_url: "https://www.airbnb.com/rooms/1176693"
name: "BEST REVIEWS★BEST MALLS★SAFE STAY★DIMSUM★CWB★MTR"
summary: "Location is OUTSTANDING as MTR Causeway BAY is within 1 min walk. You ..."
space: "FACTS - located in the heart of one of Hong Kong's most bustling place..."
description: "Location is OUTSTANDING as MTR Causeway BAY is within 1 min walk. You ..."
neighborhood_overview: "This enigmatic city of skyscrapers, ancient traditions and heavenly fo..."
notes: "Please note that not every neighbor is happy about the New Sharing Eco..."
transit: "HOW TO GET TO CAUSEWAY BAY: By Public transport: MTR: Island Line: Tin..."
access: "The entire apartment with all amenities is accessible for my guests. A..."
interaction: "I will respond to your questions within very short notice. My reservat..."
house_rules: "Early check-in and late check-out is possible for an additional fee. P..."
property_type: "Apartment"
room_type: "Entire home/apt"
bed_type: "Real Bed"
minimum_nights: "2"
maximum_nights: "1124"
cancellation_policy: "strict_14_with_grace_period"
last_scraped: 2019-03-11T04:00:00.000+00:00
```

```
use sample_airbnb
db.listingsAndReviews.find({ "bedrooms": {"$gte": 3}, "price": {"$gte": 1400, "$lte": 1500}})
```

sample_airbnb.listingsAndReviews

5.6k 4
DOCUMENTS INDEXES

Documents Aggregations Schema Explain Plan Indexes Validation

Filter ⚙️ ⏳ { "amenities": { "\$all": ["Wifi", "Kitchen"]}}}

Reset

Find

More Options ▶

⬇ ADD DATA ⏪ EXPORT COLLECTION

1 - 20 of 4767 ⏪ ⏩ ⏴ ⏵ ⏷ ⏸ ⏹

```
_id: "10047964"
listing_url: "https://www.airbnb.com/rooms/10047964"
name: "Charming Flat in Downtown Moda"
summary: "Fully furnished 3+1 flat decorated with vintage style. Located at the..."
space: "The apartment is composed of 1 big bedroom with double sized bed, a gu..."
description: "Fully furnished 3+1 flat decorated with vintage style. Located at the..."
neighborhood_overview: "With its diversity Moda- Kadikoy is one of the most colorfull neighbou..."
notes: ""
transit: ""
access: ""
interaction: ""
house_rules: "Be and feel like your own home, with total respect and love..this woul..."
property_type: "House"
room_type: "Entire home/apt"
bed_type: "Real Bed"
minimum_nights: "2"
maximum_nights: "1125"
cancellation_policy: "flexible"
last_scraped: 2019-02-18T05:00:00.000+00:00
calendar_last_scraped: 2019-02-18T05:00:00.000+00:00
first_review: 2016-04-02T04:00:00.000+00:00
last_review: 2016-04-02T04:00:00.000+00:00
accommodates: 6
```

```
use sample_airbnb
db.listingsAndReviews.find({
  "amenities": { "$all": ["Wifi", "Kitchen"]}})
```

sample_airbnb.listingsAndReviews

5.6k DOCUMENTS **4** INDEXES

Documents Aggregations Schema Explain Plan Indexes Validation

Filter !": 3}, "price": {"\$gte": 1400, "\$lte": 1500}, "amenities": { "\$all": ["Wifi", "Kitchen"] }}] Reset Find </> More Options ▾

 ADD DATA ▾

1 - 4 of 4      | 

```
_id: "1176693"
listing_url: "https://www.airbnb.com/rooms/1176693"
name: "BEST REVIEWS★BEST MALLS★SAFE STAY★DIMSUM★CWB★MTR"
summary: "Location is OUTSTANDING as MTR Causeway BAY is within 1 min walk. You ..."
space: "FACTS - located in the heart of one of Hong Kong's most bustling place..."
description: "Location is OUTSTANDING as MTR Causeway BAY is within 1 min walk. You ..."
neighborhood_overview: "This enigmatic city of skyscrapers, ancient traditions and heavenly fo..."
notes: "Please note that not every neighbor is happy about the New Sharing Eco..."
transit: "HOW TO GET TO CAUSEWAY BAY: By Public transport: MTR: Island Line: Tin..."
access: "The entire apartment with all amenities is accessible for my guests. A..."
interaction: "I will respond to your questions within very short notice. My reservat..."
house_rules: "Early check-in and late check-out is possible for an additional fee. P..."
property_type: "Apartment"
room_type: "Entire home/apt"
bed_type: "Real Bed"
minimum_nights: "2"
maximum_nights: "1124"
cancellation_policy: "strict_14_with_grace_period"
last_scraped: 2019-03-11T04:00:00.000+00:00
calendar_last_scraped: 2019-03-11T04:00:00.000+00:00
first_review: 2013-07-04T04:00:00.000+00:00
last_review: 2018-07-29T04:00:00.000+00:00
accommodates: 8
```

```
use sample_airbnb
db.listingsAndReviews.find({"address.market": "Hong Kong", "property_type": "Apartment", "bedrooms": {"$gte": 3}, "price": {"$gte": 1400, "$lte": 1500}, "amenities": { "$all": ["Wifi", "Kitchen"]}})
```

[Documents](#) [Aggregations](#) [Schema](#) [Explain Plan](#) [Indexes](#) [Validation](#)[Filter](#) [Reset](#) [Explain](#) [More Options](#)[VIEW](#) [VISUAL TREE](#) [RAW JSON](#)

Query Performance Summary [Learn more](#)

Documents Returned: **4**Index Keys Examined: **3626**Documents Examined: **3626**Actual Query Execution Time (ms): **14**Sorted in Memory: **no**Query used the following index: **PROPERTY_TYPE ↑ ROOM_TYPE ↑ BEDS ↑**

```
use sample_airbnb
db.listingsAndReviews.find({"address.market": "Hong Kong", "property_type" : "Apartment", "bedrooms": {"$gte": 3}, "price": {"$gte": 1400, "$lte": 1500}, "amenities": { "$all": ["Wifi", "Kitchen"]}}).explain()
```



Exercise Four

Indexes and Explain Plan

At the end of this, you should be able to...

Identify how a query was run using the Explain Plan

Create an index on a collection

See how an index affects query plans



View Winning Plan

```
> db.movies.  
find({ 'scores.score' : { $lt: 60 } }).explain()
```



Execution Info

```
{  
  queryPlanner: {  
    plannerVersion: 1,  
    namespace: 'sample_training.grades',  
    indexFilterSet: false,  
    parsedQuery: { 'scores.score': { '$lt': 60 } },  
    winningPlan: {  
      stage: 'COLLSCAN',  
      filter: { 'scores.type': { '$lt': 60 } },  
      direction: 'forward'  
    },  
    rejectedPlans: []  
  }  
}
```



Let's run a simple Equality, Sort, Range query

```
use sample_mflix
db.movies.find(
{
    "cast": "Bill Murray",
    "year": {$gte: 2000}
}.sort(
    {"title": 1}
)
```

```
_id: ObjectId("573a13c2f29313caabd668e0")
fullplot: "A celebrated military contractor returns to the site of his greatest c..."
> imdb: Object
year: 2015
plot: "A celebrated military contractor returns to the site of his greatest c..."
> genres: Array
rated: "PG-13"
metacritic: 40
title: "Aloha"
lastupdated: "2015-09-02 00:02:30.273000000"
> languages: Array
> writers: Array
type: "movie"
> tomatoes: Object
poster: "https://m.media-amazon.com/images/M/MV5BMTg4Mjc0NTE1NV5BMl5BanBnXkFtZT... "
num_mflix_comments: 3
released: 2015-05-29T00:00:00.000+00:00
> awards: Object
> countries: Array
< cast: Array
    0: "Bradley Cooper"
    1: "Emma Stone"
    2: "Rachel McAdams"
    3: "Bill Murray"
> directors: Array
runtime: 105
```



What results are you seeing from the Explain Plan?

The screenshot shows the MongoDB Compass interface. On the left, the sidebar displays the database structure with the 'sample_mflix' database selected. The main window shows the 'sample_mflix.movies' collection with an 'Explain Plan' tab highlighted. The query details pane shows the following MongoDB query:

```
use sample_mflix
db.movies.find(
  {
    "cast": "Bill Murray",
    "year": {$gte: 2000}
  }
).sort(
  {"title": 1}
).explain("executionStats")
```

The 'Explain Plan' tab displays the execution plan with the following stages:

- FILTER**: `{cast:"Bill Murray", year:{$gt:2000}}`
- PROJECT**
- SORT**: `{title:1}`
- COLLATION**

On the right, the results pane shows the execution statistics:

- Actual Query Execution Time (ms): 53
- Sorted in Memory: yes
- No index available for this query.

A circular progress bar at the bottom indicates a duration of 16 ms.

Local

> 9 DBS 21 COLLECTIONS C
☆ FAVORITE

Filter your data

> admin
> config
> local
> sample_geospatial
✓ sample_mflix
comments

movies

> sample_restaurants
> sample_supplies
> sample_training
> sample_weatherdata

sample_mflix.movies Explain Plan

sample_mflix.movies

Documents Aggregations Schema Explain Plan Indexes Validation

FILTER `{cast:"Bill Murray", year:{$gt:2000}}` OPTIONS EXPLAIN RESET ...

PROJECT

SORT `{title:1}` MAXTITEMS 5000

COLLATION

SKIP 0 LIMIT 0

VIEW DETAILS AS VISUAL TREE RAW JSON

Query Performance Summary

Documents Returned: 11 Actual Query Execution Time (ms): 53

Index Keys Examined: 0 Sorted in Memory: yes

Documents Examined: 23531 ▲ No index available for this query.

SORT nReturned: 11 Execution Time: 16 ms

DETAILS

Create an index for cast, title, year

The screenshot shows the MongoDB Compass interface with the following details:

- Top Bar:** Shows the connection name "SA-NA-NYC-Workshop", status "ACTIVE", and a green "OK" button.
- Left Sidebar:** Under "DATA STORAGE", the "Clusters" tab is selected, showing a cluster named "demo". Other options include "Triggers", "Data Lake", "SECURITY", and "Database Access".
- Middle Panel:** A modal window titled "Create Index" is open. It contains:
 - TIP:** A box with text about index creation best practices and a link to the "Index Strategies Documentation".
 - COLLECTION:** "sample_mflix.movies"
 - FIELDS:** A list of fields with the value "1":
 - "cast":1
 - "title":1
 - "year":1
 - OPTIONS:** A section for additional index options, currently empty.
- Bottom Panel:** A terminal window showing the MongoDB shell command to create the index:

```
use sample_mflix
db.movies.createIndex({
  "cast":1,
  "title":1,
  "year":1
})
```
- Right Side:** A summary card for the cluster "demo" with details: VERSION 4.4.0, REGION M10, and a "CREATE INDEX" button.

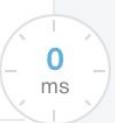
What's good and bad about the Equality, Range, Sort index we just created?

FILTER `{"cast": "Bill Murray", "year":{$gte: 2000}}` ▶ OPTIONS EXPLAIN

VIEW DETAILS AS **VISUAL TREE** RAW JSON

Query Performance Summary

Documents Returned: 12	Actual Query Execution Time (ms): 0
Index Keys Examined: 21	Sorted in Memory: no
Documents Examined: 12	Query used the following index: cast ↑ title ↑ year ↑

FETCH
nReturned: **12** Execution Time:  0 ms

DETAILS

What do you think will happen if we return only the title?

The screenshot shows the MongoDB Compass interface. On the left, the sidebar lists databases and collections, with 'sample_mflix' selected. The main area displays the 'sample_mflix.movies' collection. The 'Explain Plan' tab is active, highlighted with a green border. Below it, the query details are shown:

```
use sample_mflix
db.movies.find(
  {
    "cast": "Bill Murray",
    "year": {$gte: 2000}
  },
  {"_id": 0, "title": 1}
).explain("executionStats")
```

The Explain Plan shows the following stages:

- FILTER**: `{cast:"Bill Murray", year:{$gt:2000}}`
- PROJECT**
- SORT**: `{title:1}`

Query options include:

- MAXTIMEMS**: 5000
- SKIP**: 0
- LIMIT**: 0

Execution statistics on the right side of the interface indicate:

- Actual Query Execution Time (ms): 33
- Sorted in Memory: no
- Query used the following index:

A small circular timer icon at the bottom left shows a value of 0 ms.



Exercise Five

Aggregation Framework

*nix Pipes

```
ps ax |  
       ↗ grep mongod |  
       ↗ head 1
```

MQL Pipelines

```
match  
      ↗ group  
      ↗ sort
```



At the end of this, you should be able to...

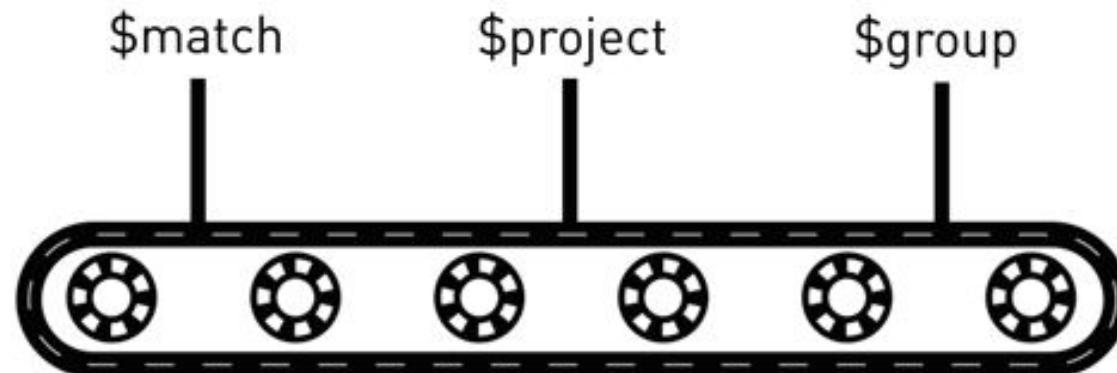
Have a basic understanding of what the aggregation framework is

Build an aggregation pipeline to perform analytics on your documents



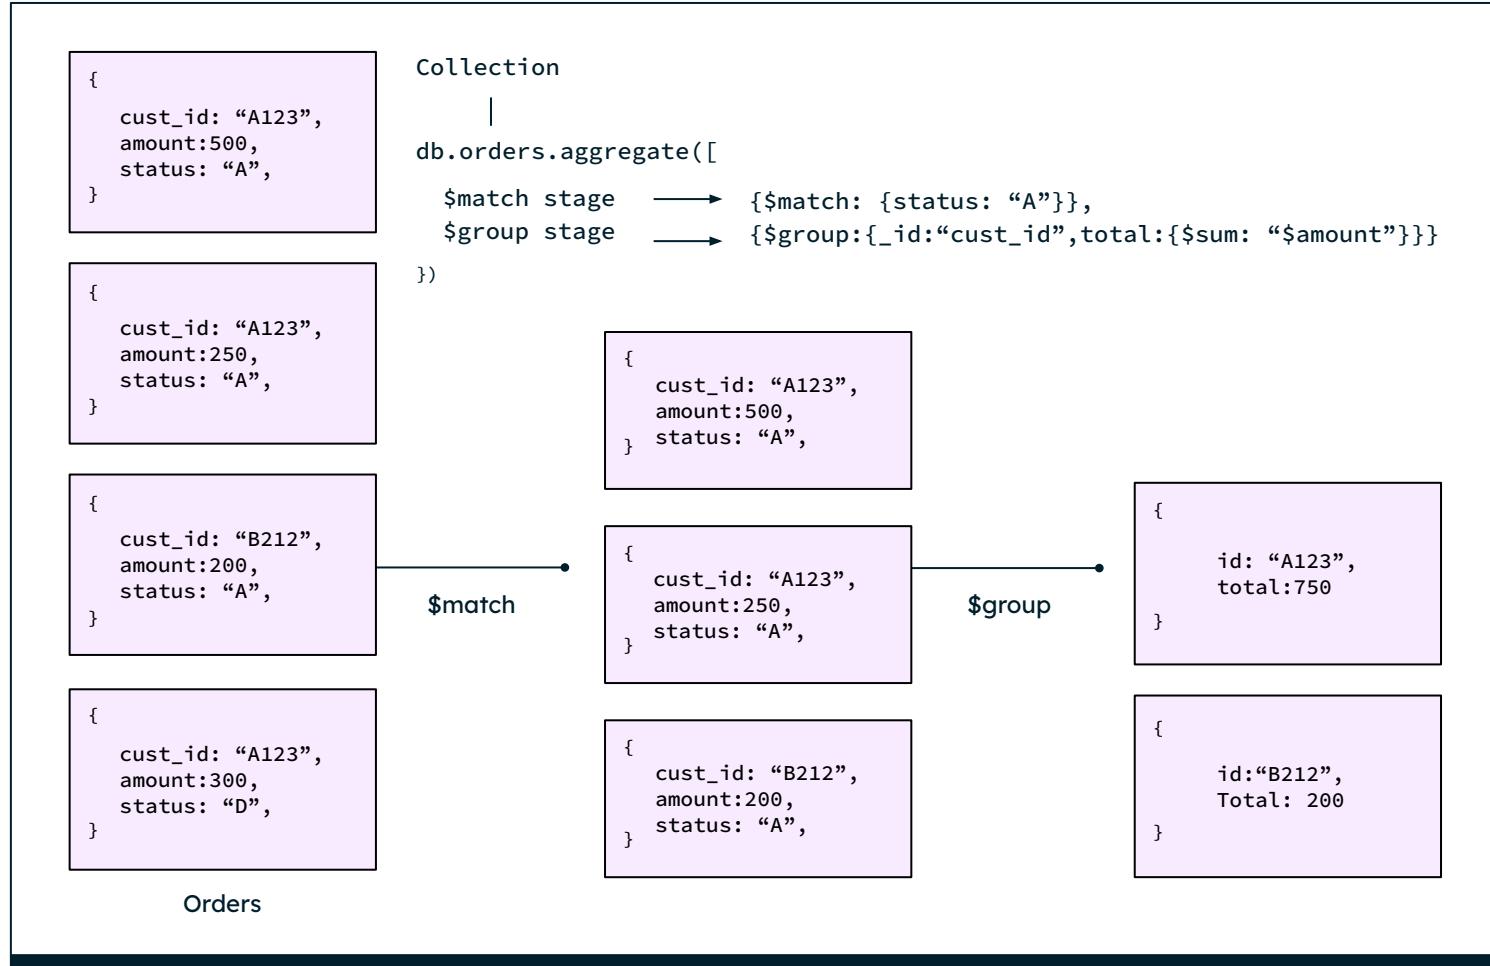
Aggregation -> Pipeline

- Each transformation is a single step known as a stage.
- Compared to one huge SQL style statement this is
 - Easier to understand
 - Easier to debug
 - Easier for MongoDB to rewrite and optimize





Aggregation pipelines





Aggregation features

A feature rich framework for data transformation and Analytics

Pipeline Stages

- \$match
- \$group
- \$facet
- \$geoNear
- \$graphLookup
- \$lookup
- \$merge
- \$project
- \$search
- \$sort
- \$setWindowFields
- \$unionWith
- \$unwind
- ...and more

Operators

- Mathematical
 - \$add, \$abs, \$subtract, \$multiply, \$divide, \$log, \$log10, \$stdDevPop, \$stdDevSam, \$avg, \$sqrt, \$pow, \$sum, \$zip, \$convert, \$round, etc.
- Array
 - \$push, \$reduce, \$reverseArray, \$addToSet, \$arrayElemAt, \$slice, etc.
- Conditionals
 - \$and, \$or, \$eq, \$lt, \$lte, \$gt, \$gte, \$cmp, \$cond, \$switch, \$in, etc.
- Temporal
 - Window Functions
 - \$dateAdd, \$dateDiff, \$dateSubtract, \$dateTrunc
 - \$dateFromParts, \$dateToParts, \$dateFromString, \$dateToString, \$dayOfMonth, \$isoWeek, \$minute, \$month, \$year, etc.
- String
 - \$toUpperCase, \$toLowerCase, \$substr, \$strcasecmp, \$concat, \$split, etc.
- Literals
 - \$exp, \$let, \$literal, \$map, \$type, etc.
- Regex
 - \$regexFnd, \$regexMatch, etc
- Trigonometry
 - \$sin, \$cos, \$degreesToRadians, etc.
- Custom Aggregation Expressions



Compared to SQL JOINs and aggregation

```
SELECT
    city,
    SUM(annual_spend) Total_Spend,
    AVG(annual_spend) Average_Spend,
    MAX(annual_spend) Max_Spend,
    COUNT(annual_spend) customers
FROM (
    SELECT t1.city, customer.annual_spend
    FROM customer
    LEFT JOIN (
        SELECT address.address_id, city.city,
               address.customer_id, address.location
        FROM address LEFT JOIN city
        ON address.city_id = city.city_id
    ) AS t1
    ON
        (customer.customer_id = t1.customer_id AND
         t1.location = "home")
) AS t2
GROUP BY city;
```

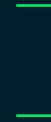
SQL queries have a nested structure

Understanding the outer layers requires understanding the inner ones

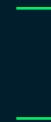
So SQL has to be read “inside-out”

Complex queries fast to create, optimize, and maintain

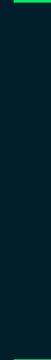
```
db.customers.aggregate([  
  {  
    $unwind: "$address",  
  },  
  {  
    $match: {"address.location": "home"}  
  },  
  {  
    $group: {  
      _id: "$address.city",  
      totalSpend: {$sum: "$annualSpend"},  
      averageSpend: {$avg: "$annualSpend"},  
      maximumSpend: {$max: "$annualSpend"},  
      customers: {$sum: 1}  
    }  
  }  
])
```



These “phases” are distinct
and easy to understand

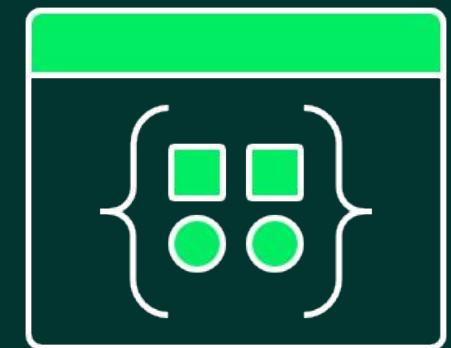


They can be thought about
in order... no nesting



MongoDB’s aggregation framework has the flexibility you need to get
value from your data, but without the complexity and fragility of SQL





Materialized Views

Faster insights on your data: pre-compute and store results of common analytics queries

With \$merge stage aggregation pipeline outputs with existing result sets to increment and enrich views

- Updated each time the pipeline is run
- Output to sharded and unsharded collections
- Define indexes on each view

With uniqueKey, control how documents are added to the view: Insert, Replace, Merge



JEOPARDY!

Metrics Collections Profiler Performance Advisor Backup Command Line

Did aggregation jeopardy

VISUALIZE YOUR DATA REFRESH

sample_mflix.movies

COLLECTION SIZE: 35.9MB TOTAL DOCUMENTS: 23539 INDEXES TOTAL SIZE: 13.79MB

Find Indexes Schema Anti-Patterns 0 Aggregation Search^{BETA}

COLLATION AUTO PREVIEW

23539 Documents in the Collection Preview of Documents in the Collection

`_id: ObjectId("573a1390f29313caabcd4135")
num_mflix_comments: 1
lastupdated: "2015-08-26 00:03:50.133000"`



Aggregation: Jeopardy

Aggregation	What's the question?
{ \$match: { genres: "Comedy" } }	



Aggregation: Jeopardy

Aggregation	What's the question?
{ \$match: { genres: "Comedy" } }	How can you find all movies with Comedy as a Genre?



Aggregation: Jeopardy

Aggregation	What's the question?
{ \$match: { genres: "Comedy" } }	How can you find all movies with Comedy as a Genre?
{ \$unwind: { path: "\$countries" } }	



Aggregation: Jeopardy

Aggregation	What's the question?
{ \$match: { genres: "Comedy" } }	How can you find all movies with Comedy as a Genre?
{ \$unwind: { path: "\$countries" } }	How can you create an individual document for each country the movie was in?



Aggregation: Jeopardy

Aggregation	What's the question?
{ \$match: {genres: "Comedy"} }	How can you find all movies with Comedy as a Genre?
{ \$unwind: {path: "\$countries"} }	How can you create an individual document for each country the movie was in?
{ \$group: { _id: "\$countries", count: {\$sum: 1} } }	



Aggregation: Jeopardy

Aggregation	What's the question?
{ \$match: {genres: "Comedy"} }	How can you find all movies with Comedy as a Genre?
{ \$unwind: {path: "\$countries"} }	How can you create an individual document for each country the movie was in?
{ \$group: { _id: "\$countries", count: {\$sum: 1} } }	How can you count all the comedies grouped by country?



Aggregation: Jeopardy

Aggregation	What's the question?
{ \$match: {genres: "Comedy"} }	How can you find all movies with Comedy as a Genre?
{ \$unwind: {path: "\$countries"} }	How can you create an individual document for each country the movie was in?
{ \$group: { _id: "\$countries", count: {\$sum: 1} } }	How can you count all the comedies grouped by country?
{ \$out: "jeopardy" }	



Aggregation: Jeopardy

Aggregation	What's the question?
{ \$match: {genres: "Comedy"} }	How can you find all movies with Comedy as a Genre?
{ \$unwind: {path: "\$countries"} }	How can you create an individual document for each country the movie was in?
{ \$group: { _id: "\$countries", count: {\$sum: 1} } }	How can you count all the comedies grouped by country?
{ \$out: "jeopardy" }	How can you output the results into a new collection, jeopardy?

DOUBLE JEOPARDY!





How many comedies does
{_id: "France"}
have based on the results in
the jeopardy collection?



TRIPLE JEOPARDY!



Exercise 2

Calculate how many Airbnb properties there are in each country, ordered by count, list the one with the largest count first. You should continue working with namespace **sample_airbnb.listingsAndReviews**.

- **Hint:** to count things, you add the explicit value 1 to an accumulator using \$sum



ANSWER: Exercise 2

Calculate how many Airbnb properties there are in each country, ordered by count, list the one with the largest count first.

```
groupfield = "$address.country"
groupstage = { $group: { _id: groupfield, count:{$sum:1}}}
sortstage = {$sort:{count:-1}}
pipe = [groupstage,sortstage]
db.listingsAndReviews.aggregate(pipe).pretty()
```

or

```
db.listingsAndReviews.aggregate([{$sortByCount:"$address.country"}]).pretty()
```

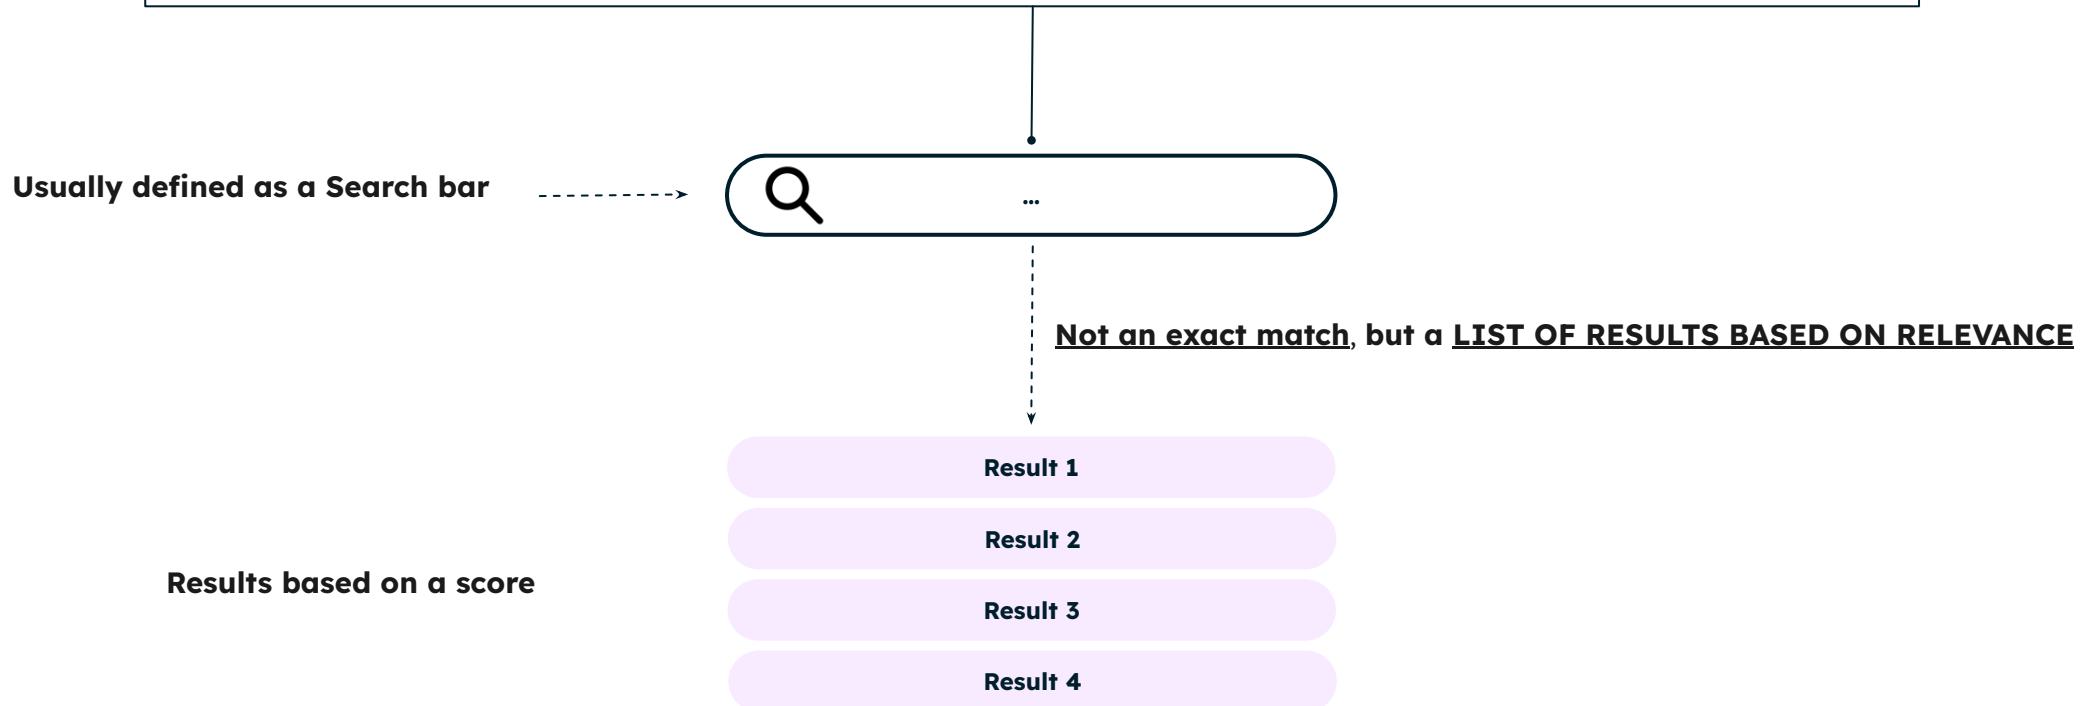


Search



What is Search?

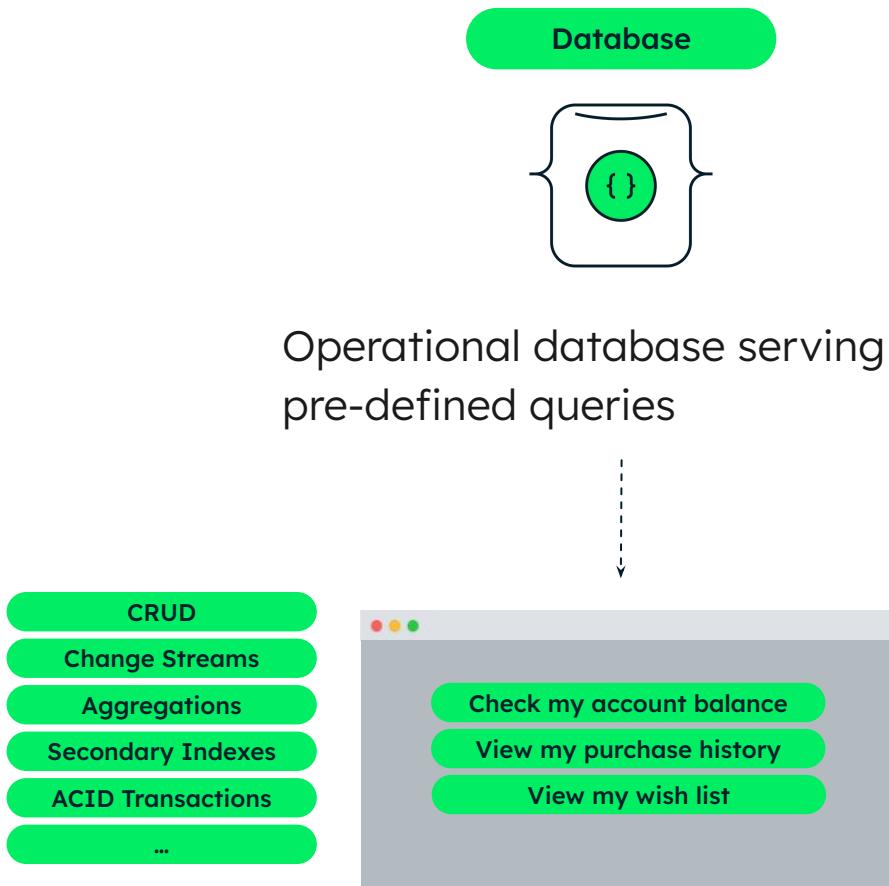
Search, or “Full-Text Search” is the ability to search across all of your data and efficiently return a list of results, ranked based on how well they matched to the search term





How developers build app search today

Stand up a database to provide the application's persistence layer



How developers build app search today



Bolt-on a search engine to power search across application data

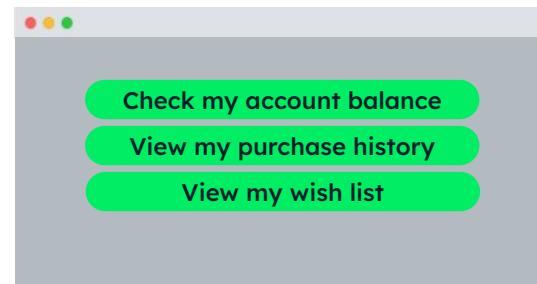
Database



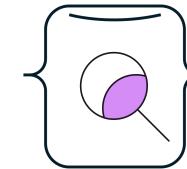
Operational database serving
pre-defined queries



- CRUD
- Change Streams
- Aggregations
- Secondary Indexes
- ACID Transactions
- ...



Search Engine

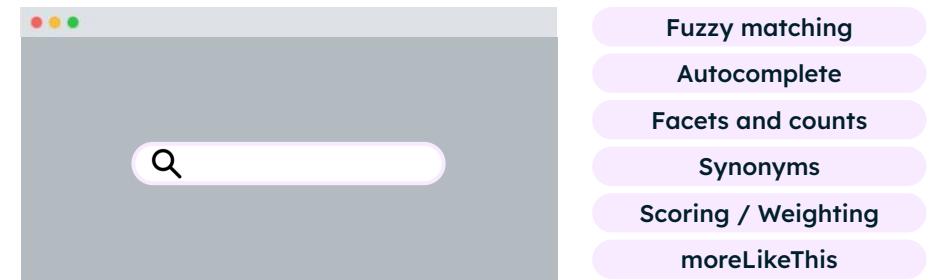


Lucene

elasticsearch

Solr

Search engine inferring intent from free
form, natural language queries

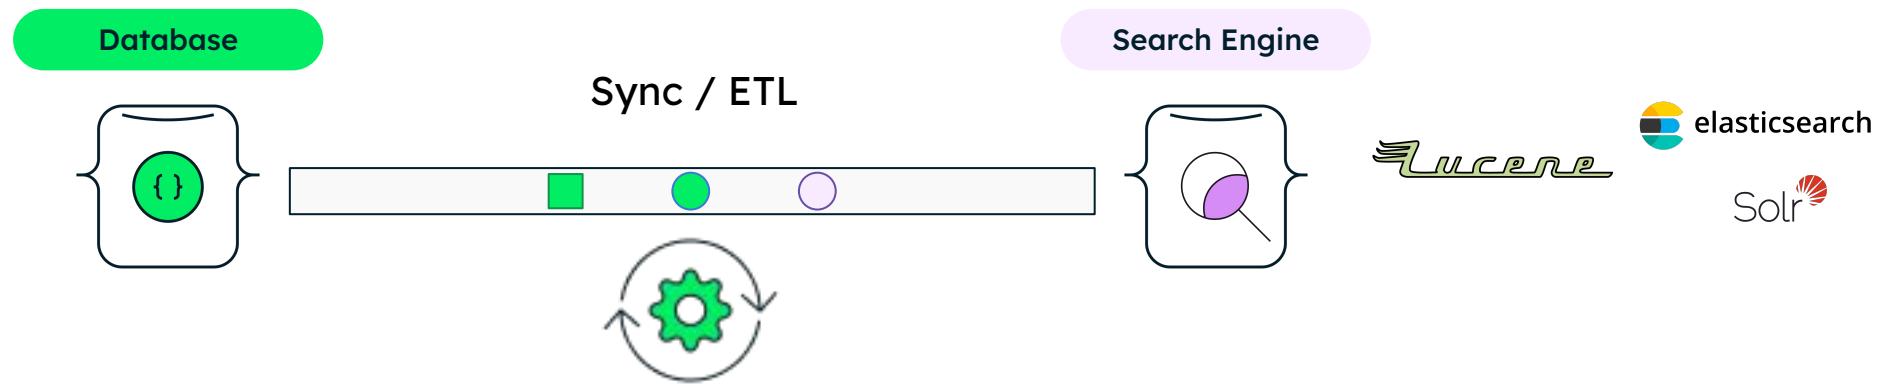


- Fuzzy matching
- Autocomplete
- Facets and counts
- Synonyms
- Scoring / Weighting
- moreLikeThis



How developers build app search today

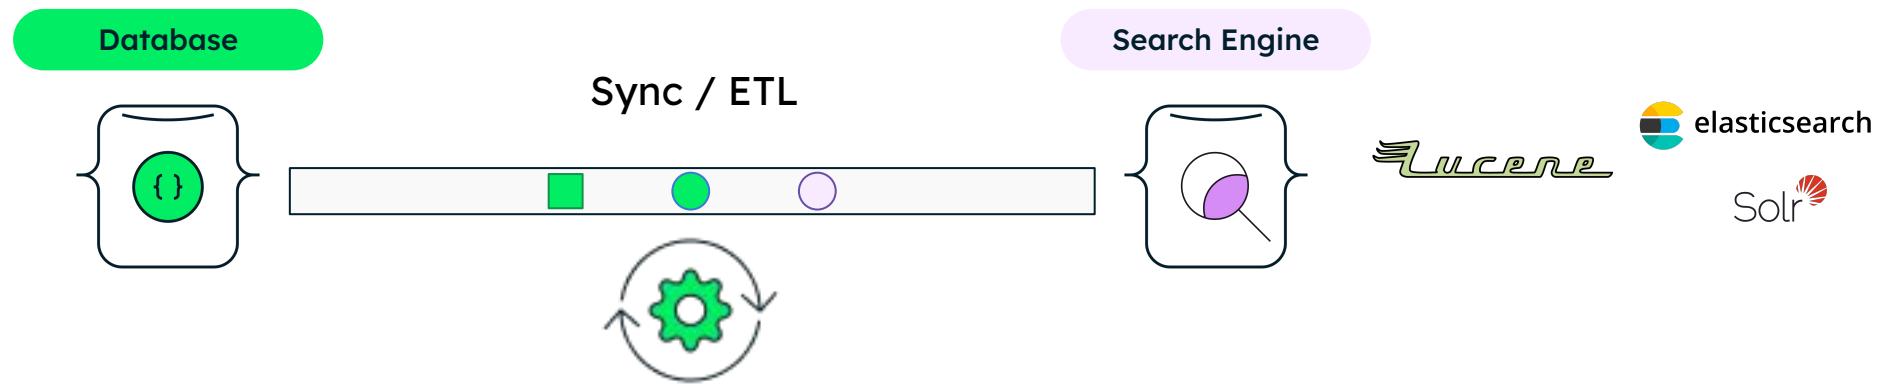
Stand up a replication mechanism to keep the systems in sync





How developers build app search today

Architectural complexity in our application stack



✖ Developer friction

Different query APIs and drivers for database and search, coordinate schema changes

✖ Pay the sync tax

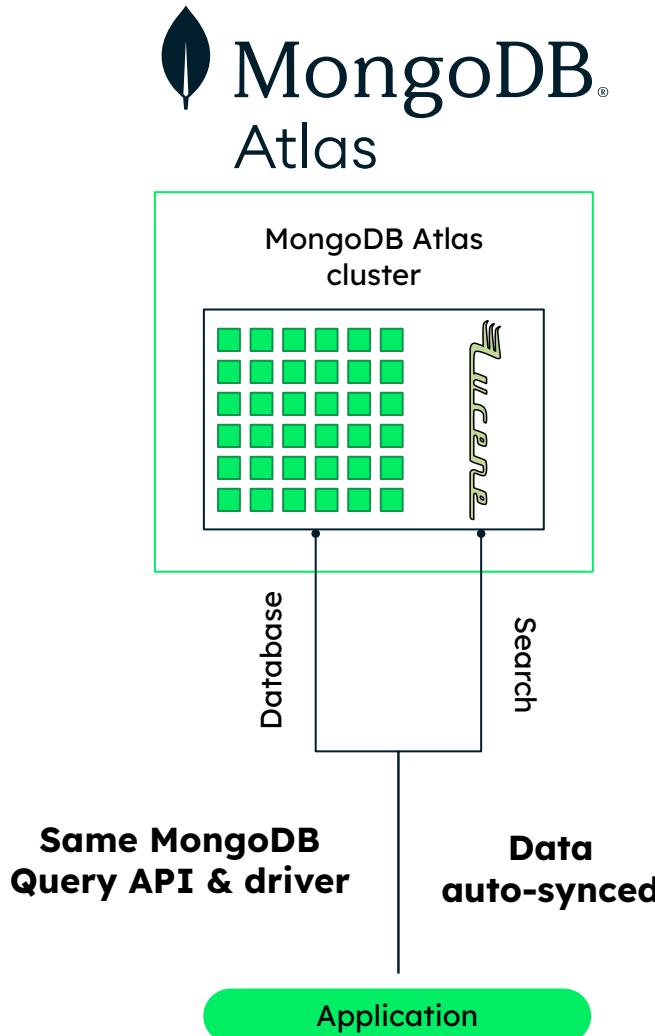
Requires its own systems and skills. Recovering sync errors can consume 10% of a developer's time

✖ Operational overhead

More to provision, secure, upgrade, patch, back up, monitor, scale, etc.



How does Atlas Search address that pain?



Embeds a fully managed Apache Lucene index right alongside the database

✓ Improved developer productivity

Build database and search features using the same query API and driver

✓ Simplified data architecture

Automatic data synchronization, even as your data and schema changes

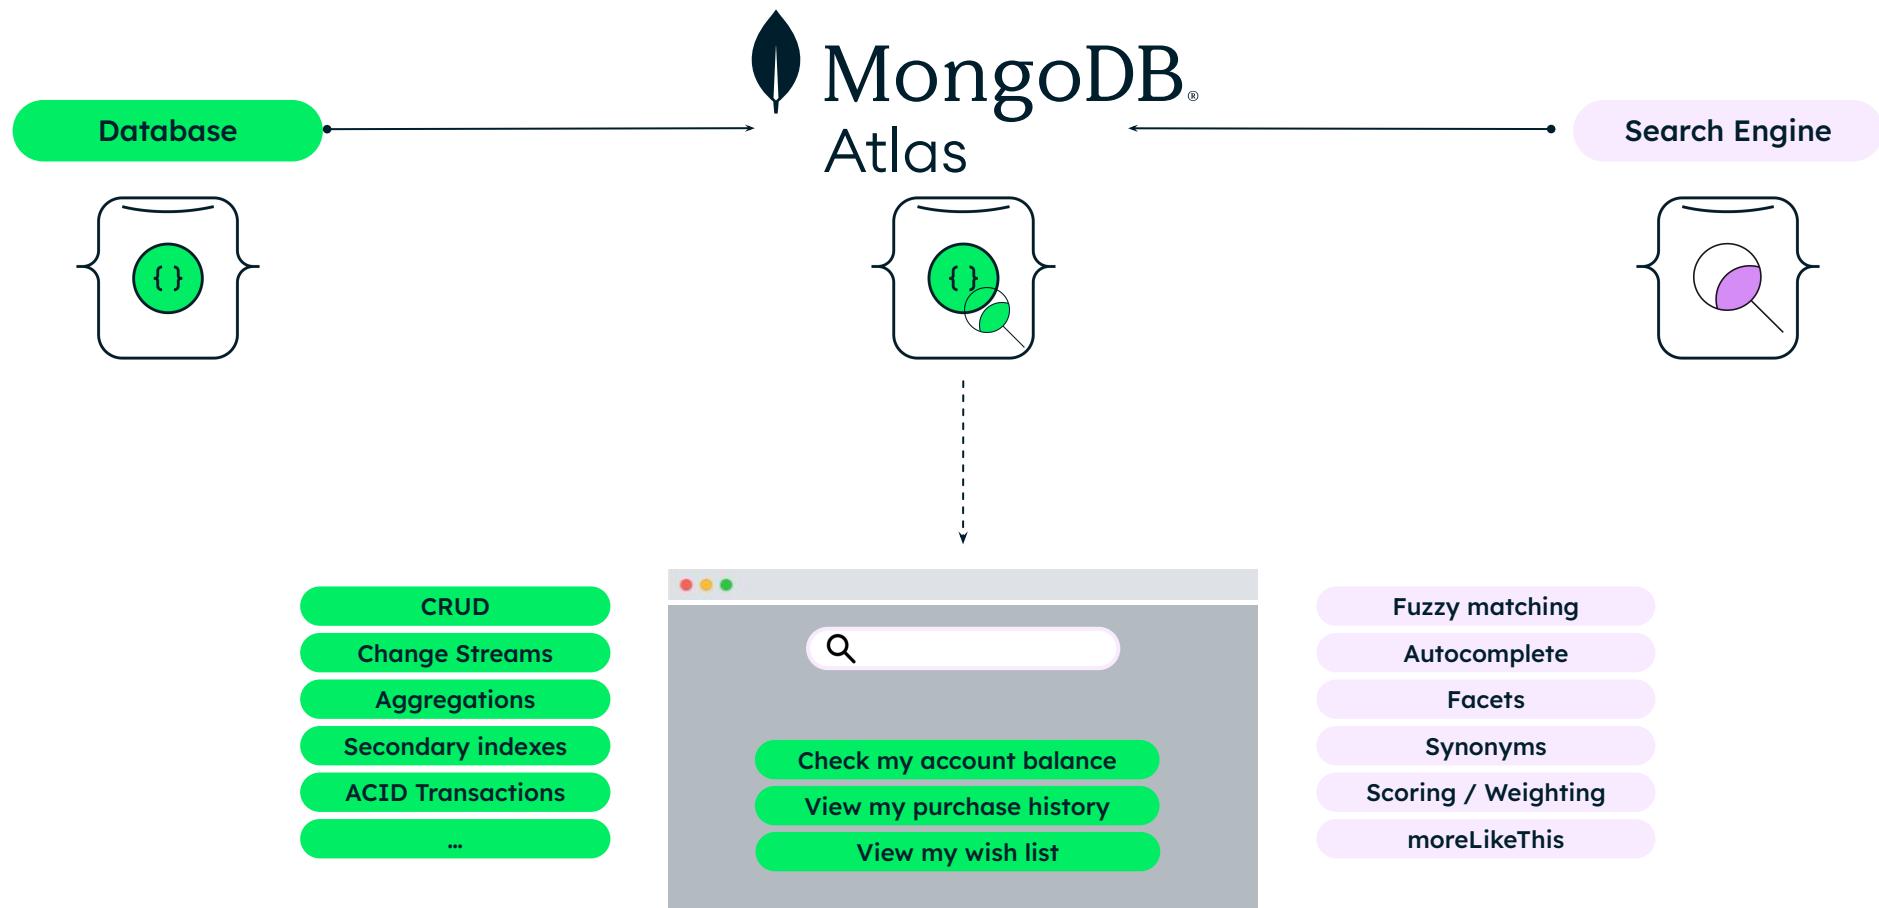
✓ Fully managed platform

Get the security, performance, and reliability of Atlas



Unifying database and search

Compress 3 systems into 1, ship search features **30%-50% faster**





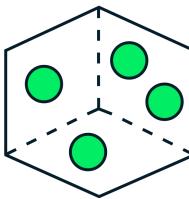
Vector Search



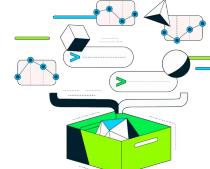
SPOTLIGHT

Atlas Vector Search

(public preview)



Build intelligent applications powered by semantic search and generative AI over any type of data



Store vector embeddings right next to your source data and metadata with the power of the document model. Vectors inserted or updated in the database are automatically synchronized to the vector index



Remove operational heavy lifting with the battle tested, fully managed Atlas platform



Why is Vector Search needed ?



- Syntactically similar sentences can have very different meanings.
- Syntactic search = matching of search terms
- Semantic search = meaning and context

Sentence#1: “**Mike Tyson knocked out Michael Spinks.**”

Sentence#2: “**Michael Spinks knocked out Mike Tyson.**” (not true btw)

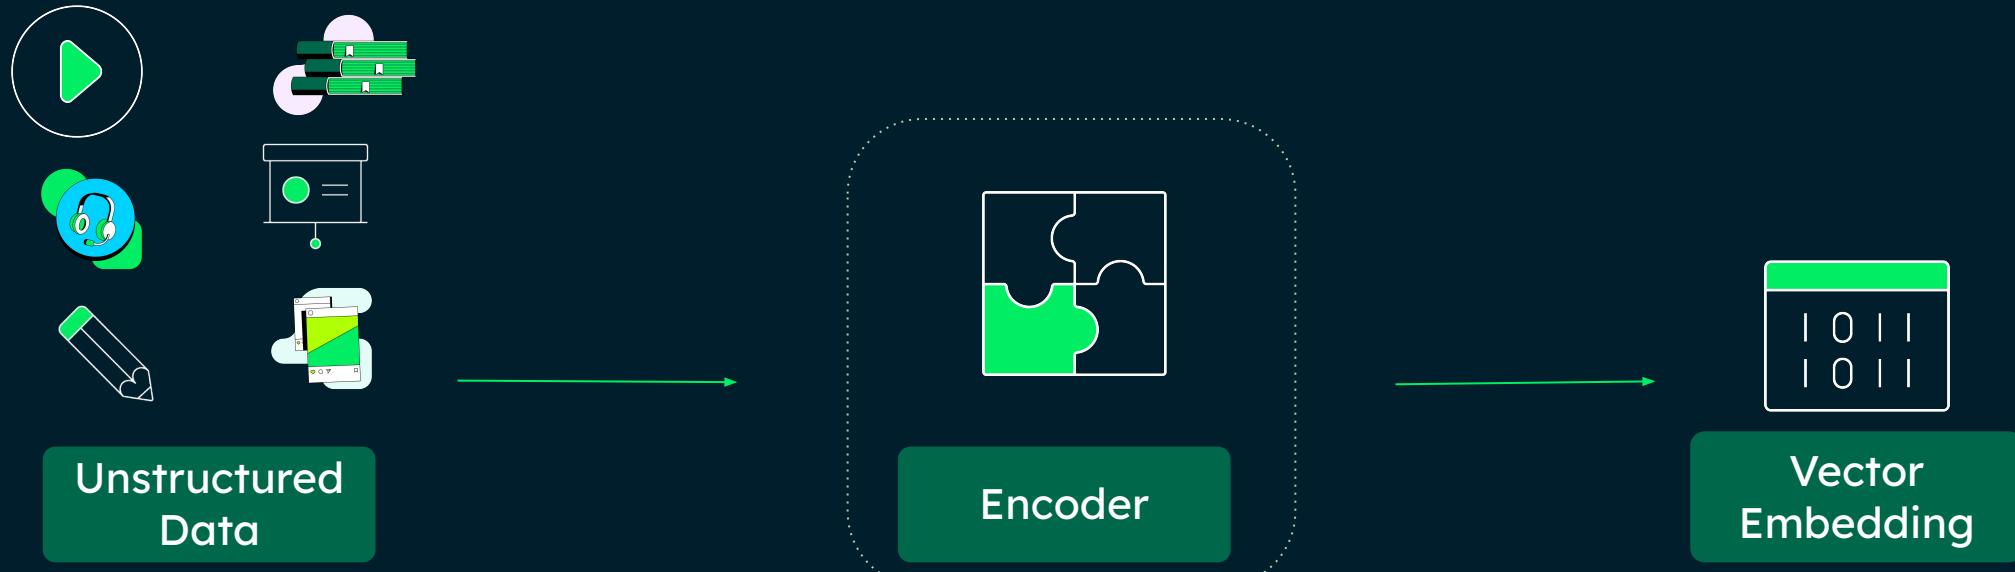


Vectors are a numeric representation
of data and related context

“a quick brown fox” = [0.743, 0.720, -0.325, 0.195, 0.835, -0.945, ...]



Vector embeddings are produced by sending data through an encoding model



E.g. OpenAI, Cohere, Anthropic,
HuggingFace, Vertex, etc.



Vector Search powers a number of key use cases

Semantic Search

Similarity Search

Recommendation Engines

Long-term memory for LLMs

Q & A Systems

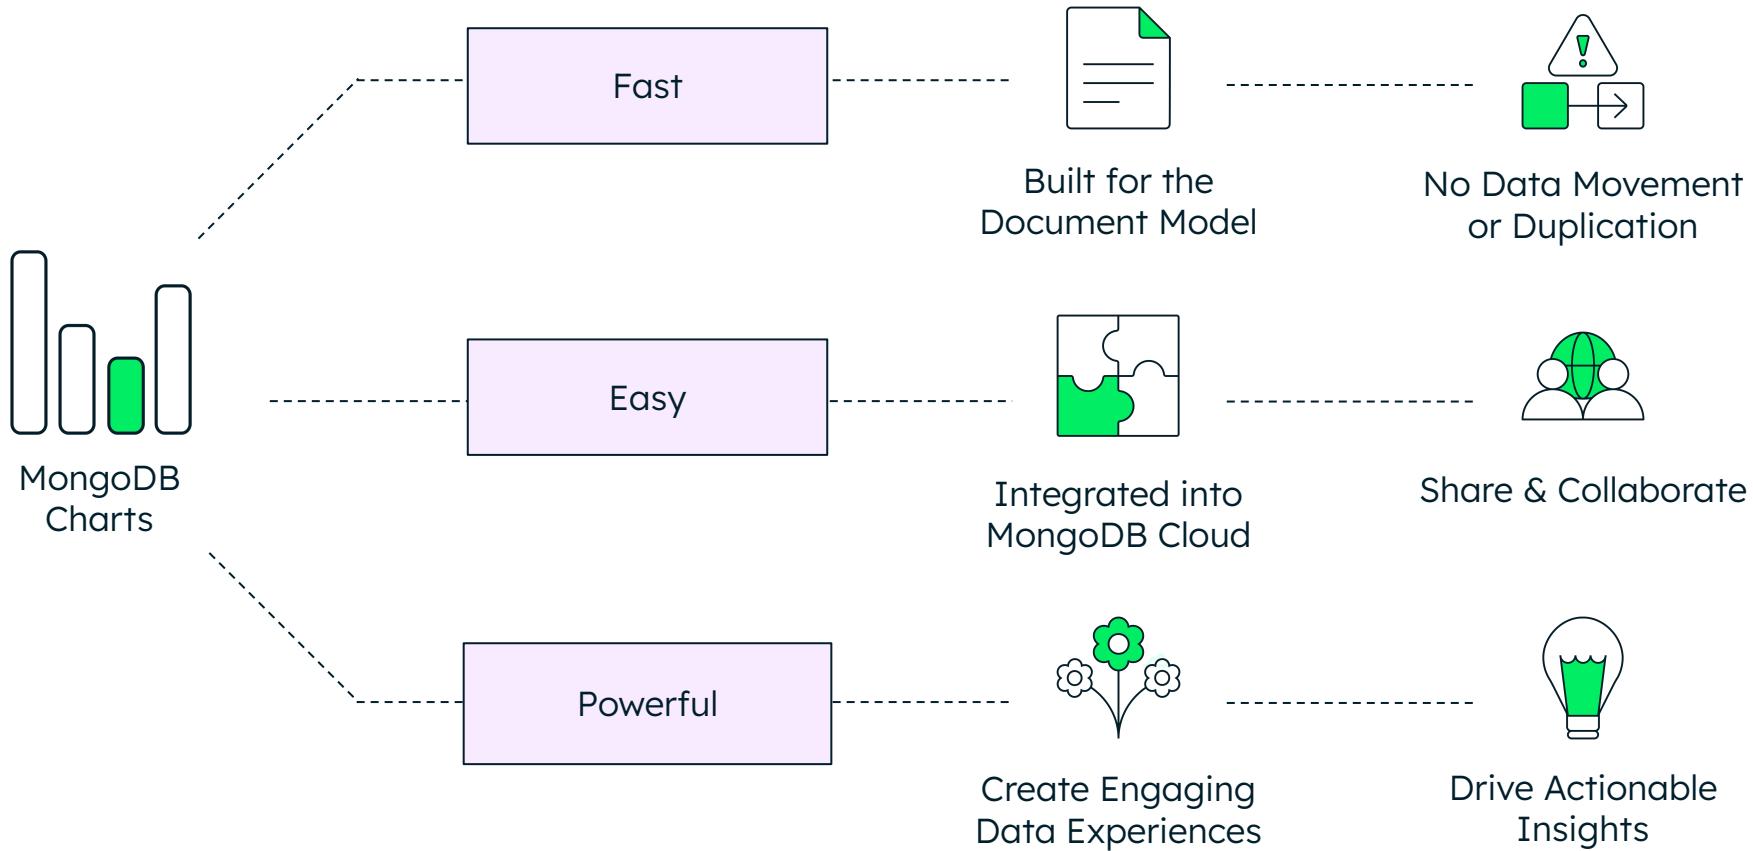
Image, Audio, Multimedia Search



MongoDB Charts

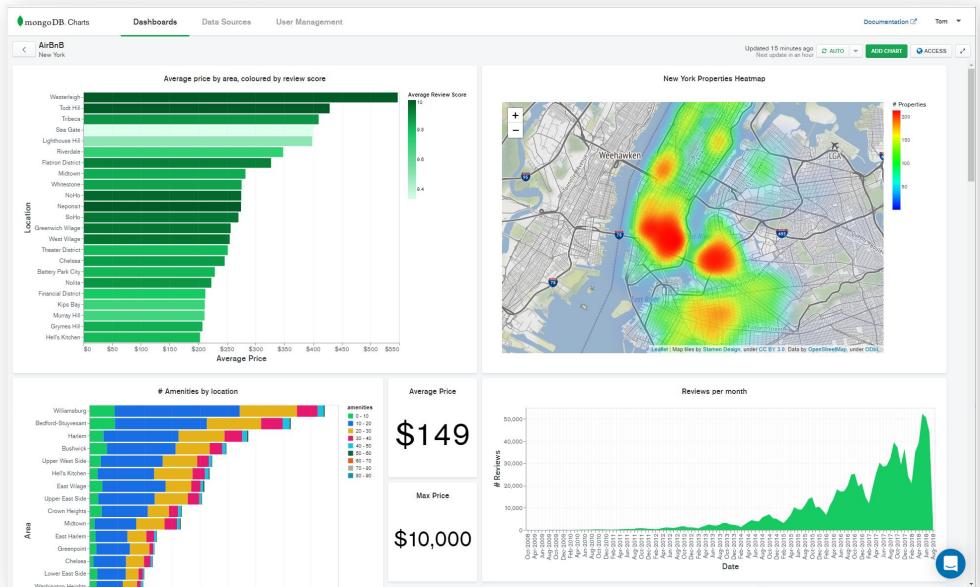


The best way to visualize MongoDB Data



Easy to get started and share

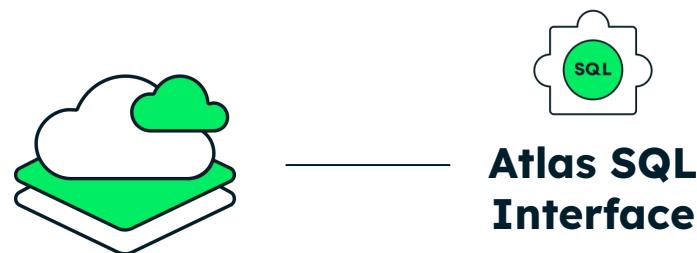
- Integrated into the MongoDB Data Application Platform, there's no setup required and total cost of ownership is reduced
- Share and collaborate with dashboards so your teams can work together or self-serve to make decisions in real-time





Atlas SQL Interface, Connectors, and Drivers

Leverage SQL tools to easily query and generate rapid insights from MongoDB Atlas



Atlas SQL Interface

Visualize data using **Custom Connectors** for BI tools



Power BI



Use **JDBC/ODBC Drivers** for other SQL connections

Bhavik

ACTIVE

Atlas

App Services

Charts



Dashboards

Data Sources

Embedding

Data Usage

Dashboards

All dashboards

Last modified date

Title or description...

Add Dashboard

Demo Dash



Cast with Most Awards, Split by Genre

inspections

2 charts

Modified 10 days ago

Charts workshop: Step 1

Create a New Dashboard

1. Click the Dashboards tab
2. Click the New/Add Dashboard button
3. Enter the Title: Movie Details
4. Click Create



Charts workshop: Step 2

Create Chart and Select Data Source

1. Click Add Chart button
2. Select the sample_mflix.movies data source under Project and demo cluster
3. Enter the Title: Movie Details



Charts workshop: Step 3

Create Chart Showing Directors with the Most Awards

1. Select Chart Type: Column / Stacked
2. X Axis: director
 - a. Sort By: Aggregated Value, Descending
 - b. Limit: 10



Charts workshop: Step 4

Create Chart Showing Directors with the Most Awards

5. Y Axis: awards.wins
 - a. Aggregate: sum
6. Series: genres
 - a. Array Reduction: Unwind Array
7. Title: Directors with Most Awards, Split by Genre
8. Click Save and Close





Thank you



MongoDB Technical Overview

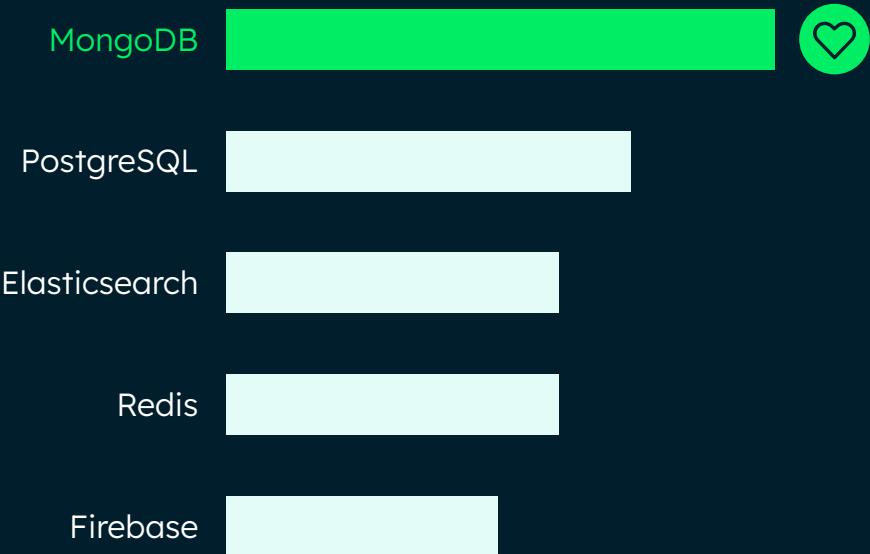


The evolution of MongoDB



Why MongoDB?





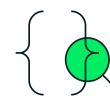
Consistently voted
“most wanted” database

STACK OVERFLOW | [DEVELOPER SURVEY](#)

Developers love using MongoDB



Intuitive and flexible
document data model



Powerful query engine and
idiomatic API to work with data



Distributed systems architecture
for resilience & scale-out



The Document Model and MongoDB Query API

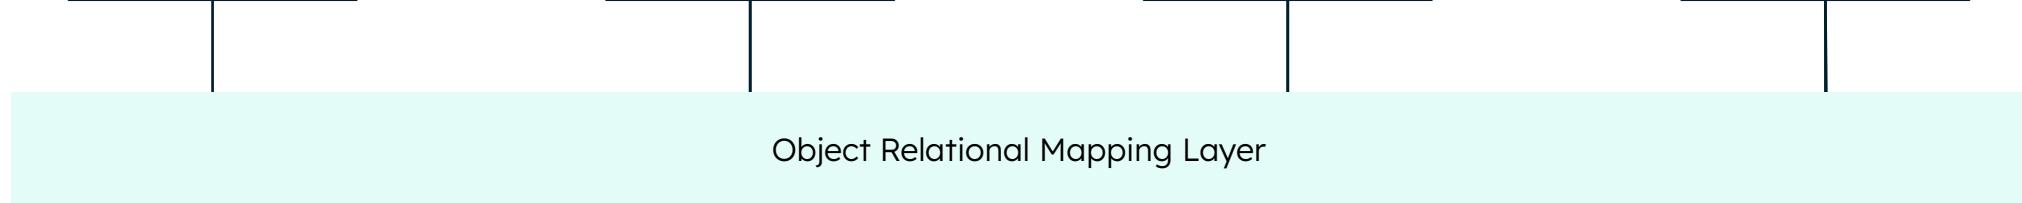
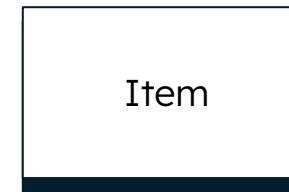
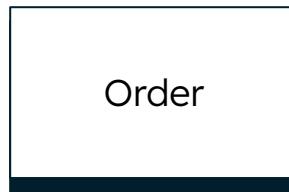
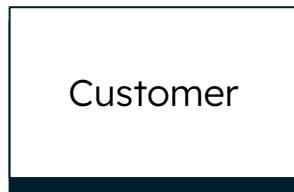
The fastest way to innovate

You probably have thousands of tables

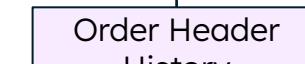
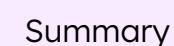
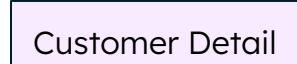


Go from this...

Objects

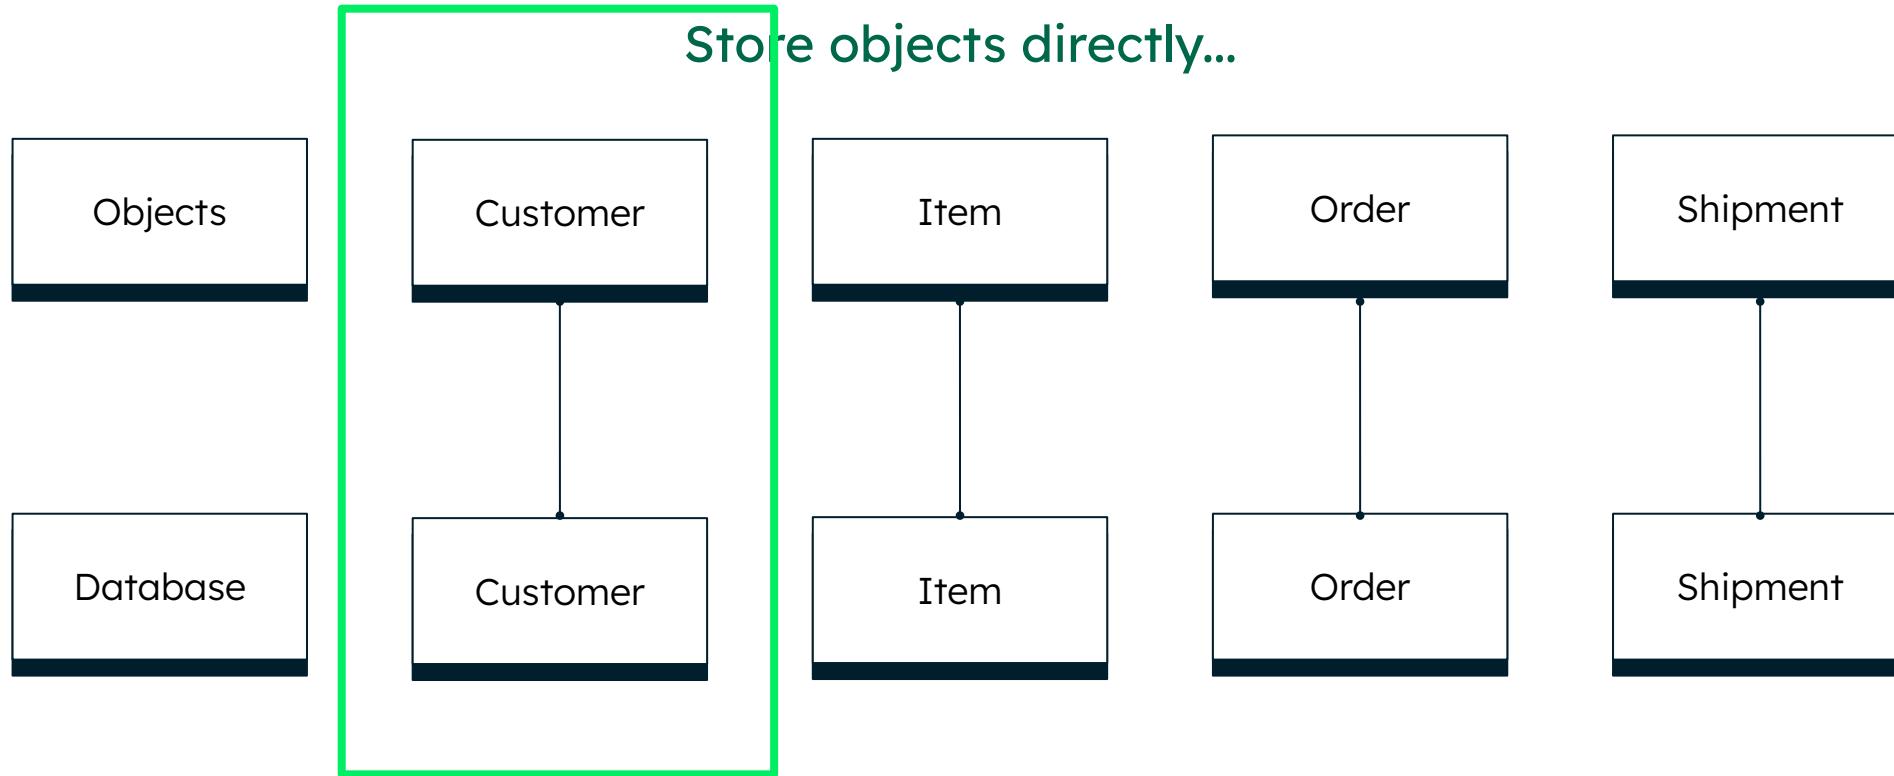


Tables





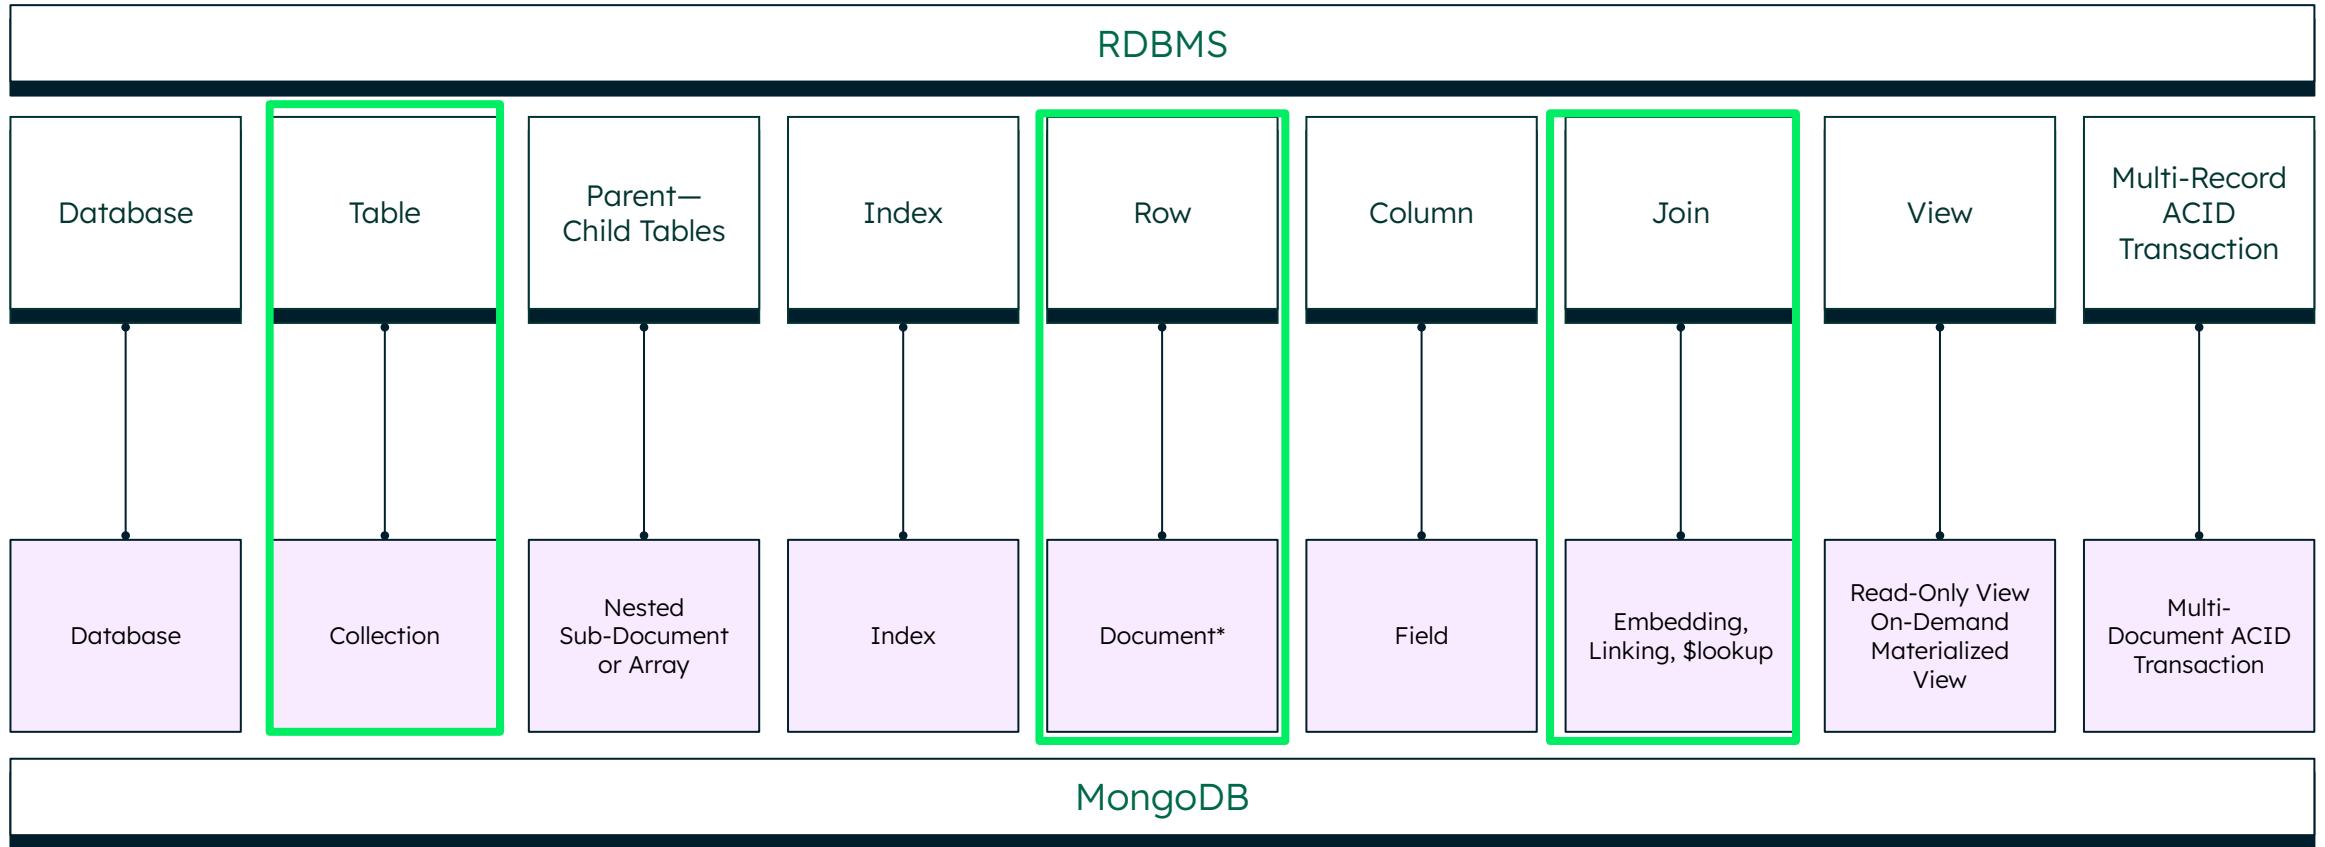
...to this





Some terminology

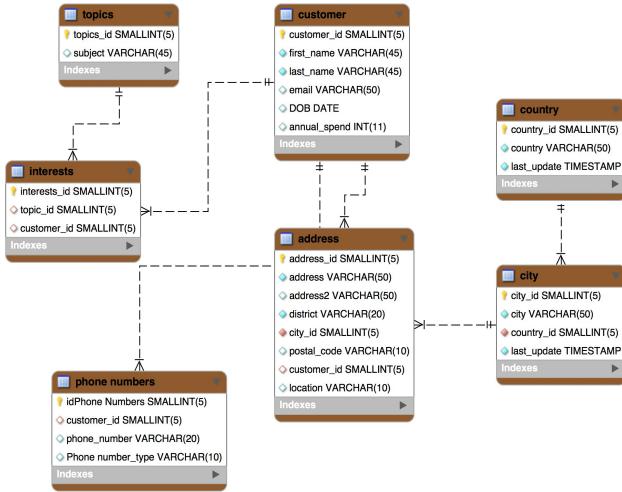
A comparison



* Proper document schema design yields more entity data per document than found in a relational database row

Contrasting data models

INTUITIVE



Tabular (Relational) Data Model

Related data split across
multiple records and tables

Document Data Model

Related data contained in a single, rich document

```
{  
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),  
  "name" : {  
    "first" : "John",  
    "last" : "Doe" },  
  "address" : [  
    { "location" : "work",  
      "address" : {  
        "street" : "16 Hatfields",  
        "city" : "London",  
        "postal_code" : "SE1 8DJ"},  
        "geo" : { "type" : "Point", "coord" : [  
          51.5065752,-0.109081]}},  
    +   {...}  
  ],  
  "dob" : ISODate("1977-04-01T05:00:00Z"),  
  "retirement_fund" : NumberDecimal("1292815.75")  
}
```



INTUITIVE

Naturally maps to objects in code

- Eliminates requirements to use ORMs
- Breaks down complex interdependencies between developer and DBAs teams

Represent data of any structure

- Polymorphic: each document can contain different fields
- Modify the schema at any time

Strongly typed for ease of processing

- Over 20 binary encoded JSON data types

Access by idiomatic drivers in all major programming language

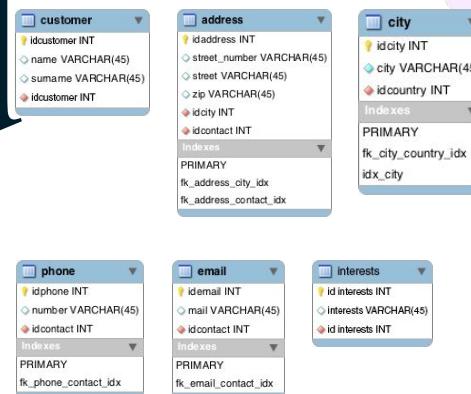
Document data model

```
{  
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),  
  "name" : {  
    "first" : "John",  
    "last" : "Doe" },  
  "address" : [  
    { "location" : "work",  
      "address" : {  
        "street" : "16 Hatfields",  
        "city" : "London",  
        "postal_code" : "SE1 8DJ"},  
        "geo" : { "type" : "Point", "coord" : [  
          51.5065752,-0.109081]}},  
    +   {...}  
  ],  
  "dob" : ISODate("1977-04-01T05:00:00Z"),  
  "retirement_fund" : NumberDecimal("1292815.75")  
}
```

Intuitive and fast

Compared to storing data across multiple tables,
a single document data structure:

- Presents a single place for the database to read and write data
- Denormalized data eliminates JOINs for most operational queries
- Simplifies query development and optimization



```
_id: 12345678
> name: Object
> address: Array
> phone: Array
email: "john.doe@mongodb.com"
dob: 1966-07-30 01:00:00:000
< interests: Array
  0: "Cycling"
  1: "IoT"
```



A row in a RDBMS table is most similar
to a _____ in MongoDB?

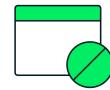
- a) Object
- b) Document
- c) Collection
- d) Record



Flexible Schema: unlocking developer velocity



Avoids need to update
ORM class mappings and
recompile programming
language classes



Schema changes don't lock
the database, or cause
performance degradation
while tables are altered



Breaks down complex
inter-group dependencies
and expensive coordination
before new code is released



Flexible: adapt to change

Add new fields dynamically at runtime

```
{  
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),  
  "name" : {  
    "first" : "John",  
    "last" : "Doe" },  
  "address" : [  
    { "location" : "work",  
      "address" : {  
        "street" : "16 Hatfields",  
        "city" : "London",  
        "postal_code" : "SE1 8DJ"},  
      "geo" : { "type" : "Point", "coord" : [  
        51.5065752,-0.109081]}},  
    + {...}  
  ],  
  "dob" : ISODate("1977-04-01T05:00:00Z"),  
  "retirement_fund" : NumberDecimal("1292815.75")  
}
```

```
{  
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),  
  "name" : {  
    "first" : "John",  
    "last" : "Doe" },  
  "address" : [  
    { "location" : "work",  
      "address" : {  
        "street" : "16 Hatfields",  
        "city" : "London",  
        "postal_code" : "SE1 8DJ"},  
      "geo" : { "type" : "Point", "coord" : [  
        51.5065752,-0.109081]}},  
    + {...}  
  ],  
  "phone" : [  
    { "location" : "work",  
      "number" : "+44-1234567890"},  
    + {...}  
  ],  
  "dob" : ISODate("1977-04-01T05:00:00Z"),  
  "retirement_fund" : NumberDecimal("1292815.75")  
}
```

Data governance



JSON Schema

Enforces strict schema structure over a complete collection for data governance & quality

- Builds on document validation introduced by restricting new content that can be added to a document
- Enforces presence, type, and values for document content, including nested array
- Simplifies application logic

Tunable

Enforce document structure, log warnings, or allow complete schema flexibility

Queryable

Identify all existing documents that do not comply



What is the statement equivalent to
ALTER TABLE in MongoDB?

- a) ALTER COLLECTION
- b) ALTER DOCUMENT
- c) ALTER TABLE
- d) None of the above



Fully featured secondary indexes—document optimized—extended beyond RDBMS experiences



Index types

Primary index

Every collection has a primary key index

Compound index

Index against multiple keys in the document

Multikey index

Index into arrays

Wildcard index

Auto-index all matching fields, sub-documents & arrays

Text indexes

Support for text searches

Geospatial indexes

2d & 2d sphere indexes for spatial geometries

Hashed indexes

Hashed based values for sharding

Index features

TTL indexes

Single field indexes, when expired delete the document

Unique indexes

Ensures value is not duplicated

Partial indexes

Expression based indexes, allowing indexes on subsets of data

Case insensitive indexes

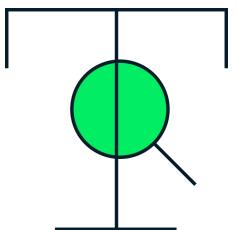
Supports text search using case insensitive search

Sparse indexes

Only index documents which have the given field

Powerful:

ATLAS Search \$SEARCH single command integrating Lucene engine





Multi-cloud global database

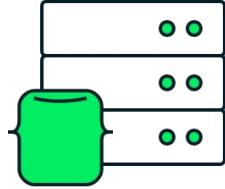
Freedom and flexibility



Broadest reach: private, hybrid, public



Laptop



Mainframe



Server



Self-managed
in the cloud



Database
as a Service

Consistent developer experience

Same codebase, same APIs, same tools, wherever you run



Don't I get this freedom with relational databases?

Fragmentation across different cloud platforms

- Different versions
- Different database features and options
- Different scaling characteristics
- Different cross-region options
- Different monitoring, alerting, backup, and recovery
- Different security controls

Architectural mis-alignment



High availability: replica sets

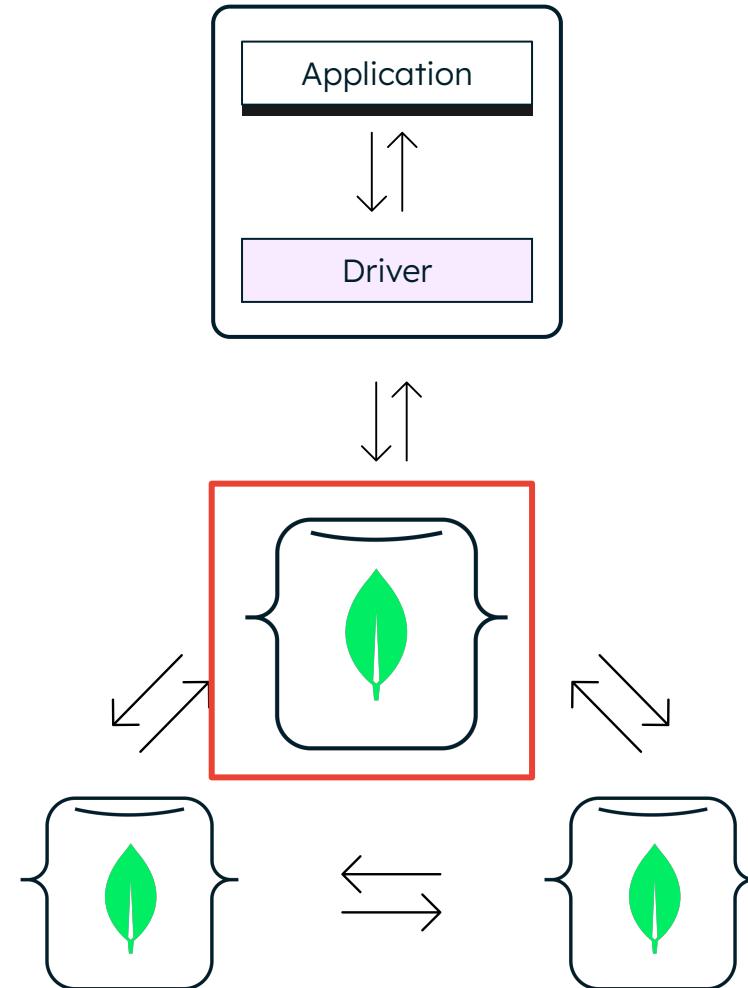
Automatic failover

- Typical failover in 5 seconds or less
- Retriable reads and writes to catch temporary exceptions

Data center aware, tunable durability, and consistency

Addresses availability considerations:

- High availability
- Disaster recovery
- Maintenance





How many hours of downtime should I plan on for maintenance?

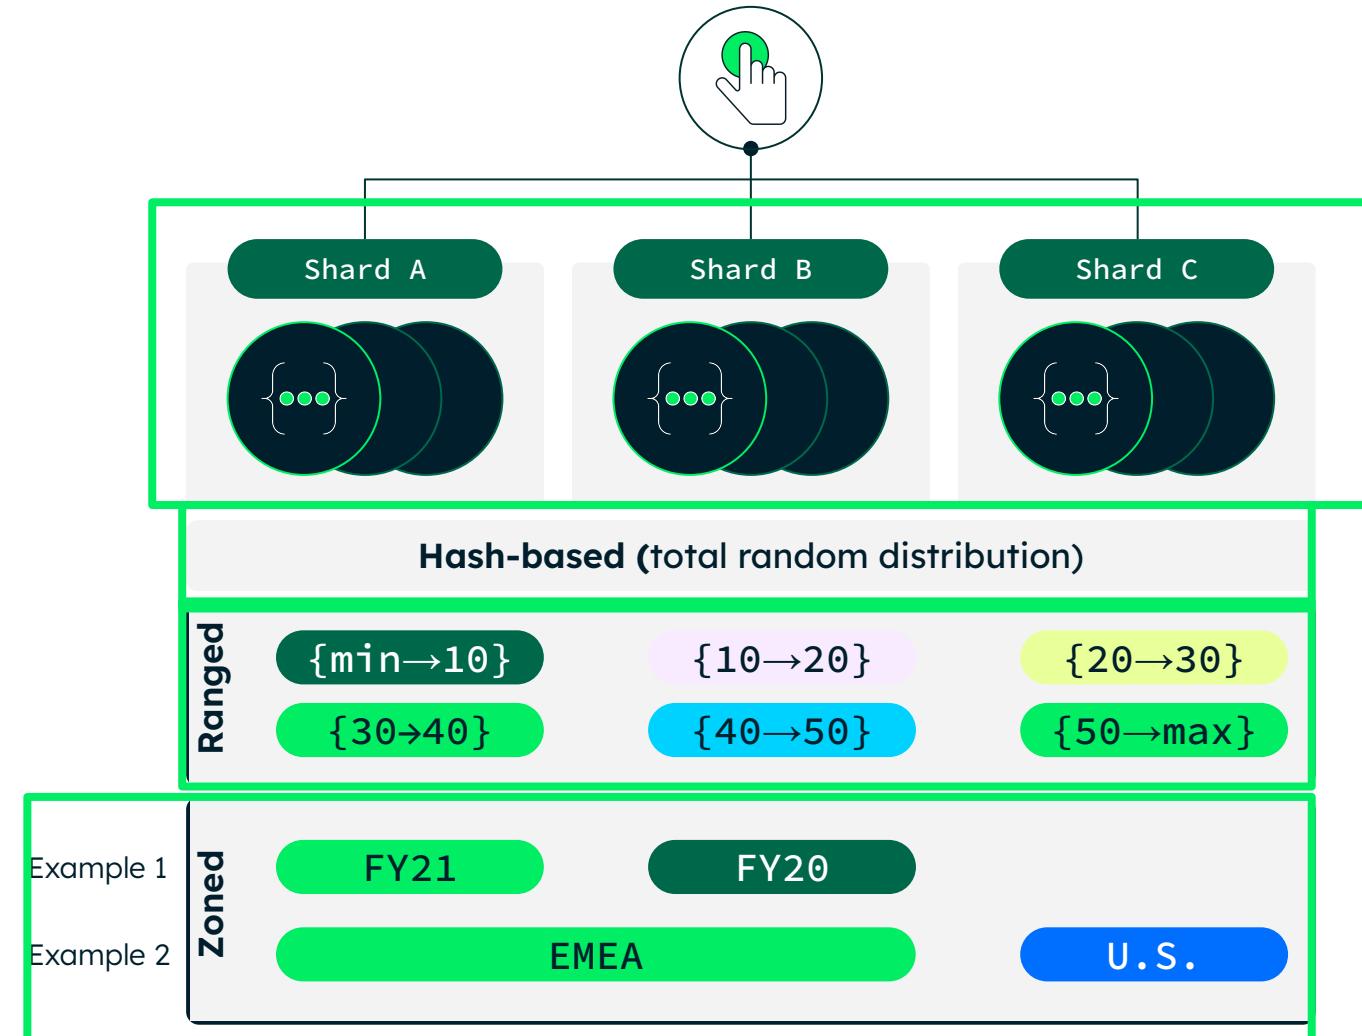
- a) 0
- b) 2
- c) 1
- d) 3



Scalability: sharding

Native-sharding for horizontal scale-out

- Automatically scale beyond the constraints of a single node
- Application transparent
- Scale and rebalance incrementally, in real time
- Rather than randomly spray data across a cluster, exposes multiple data distribution policies to optimize for query patterns and locality



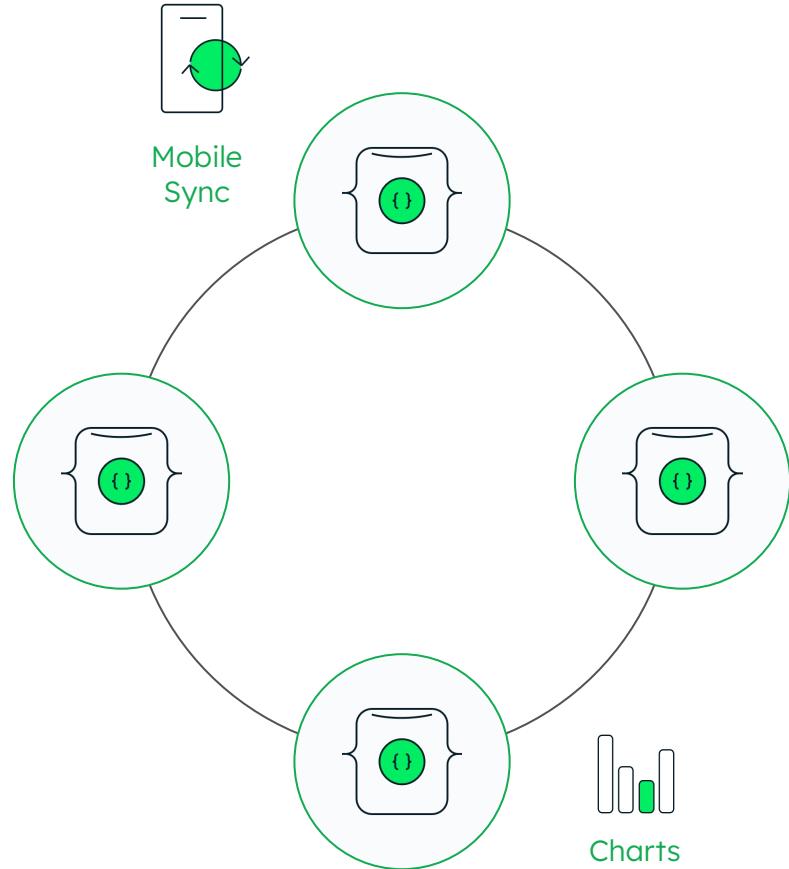


Which sharding policy would get data closest to the end user?

- a) Hashed
- b) Range
- c) Zoned



Workload isolation



Enable different workloads on the same data

- Combine operational and analytical workloads on a single data platform
- Extract live insights from real-time data to enrich applications
- One set of nodes serving operational apps, replicating to dedicated nodes serving analytics: up to 50 nodes in a single replica set
- ETL-free



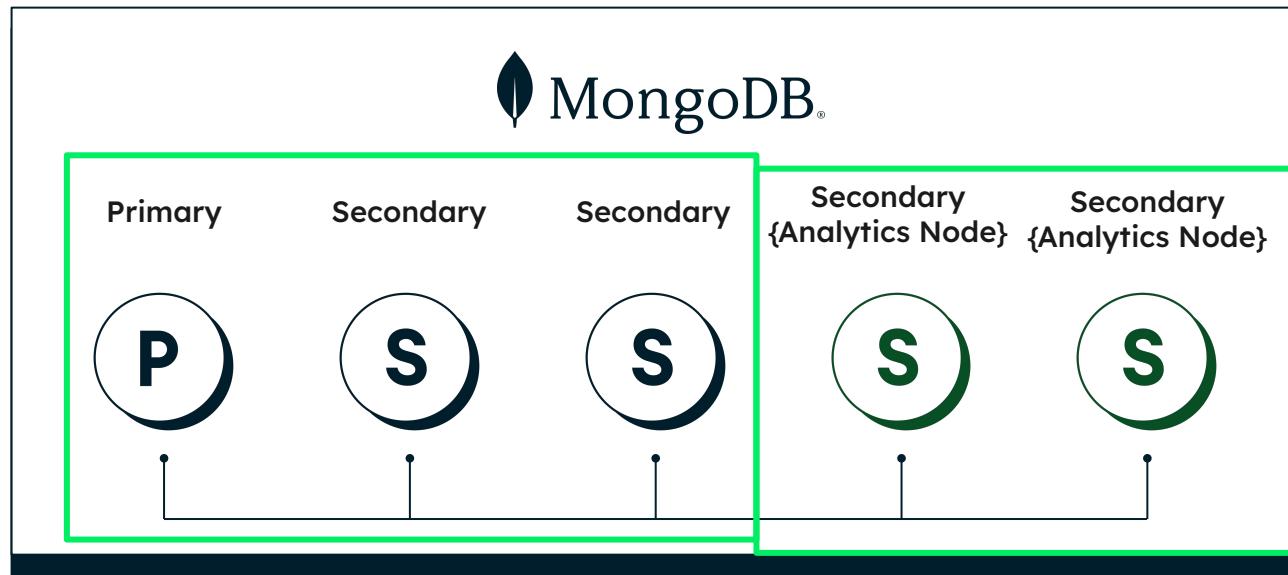
MongoDB workload isolation

Transactional Applications



Rich MongoDB Query API and distributed architecture allows you to run both Transactions and Analytics on the same cluster with no resource contention

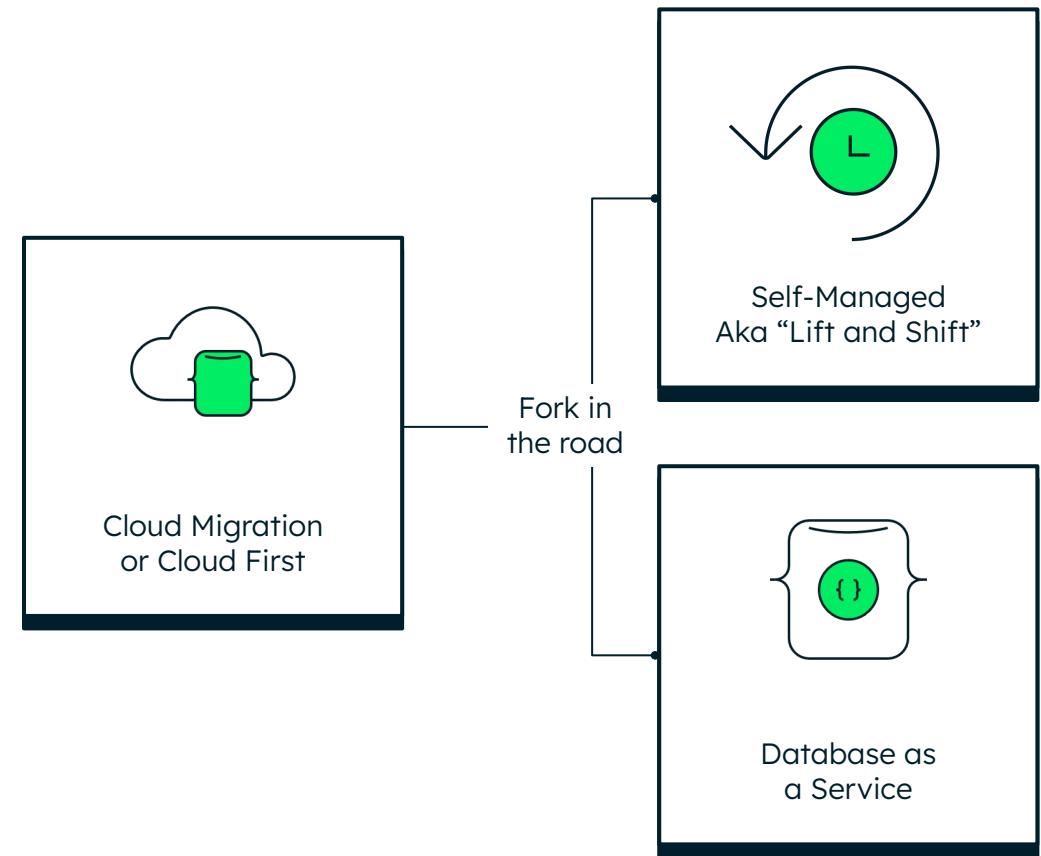
Analytics Consumers



2 choices as you move to the cloud: Self-Managed or DBaaS

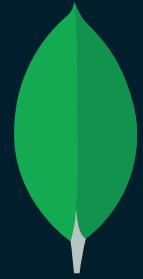
Self-managed

1. Provision instances and storage
2. Configure HA
3. Configure security
4. Configure backup/restore
5. Monitoring & alerting
6. Ongoing upgrades & maintenance



Fully-managed DBaaS

Choose instance, hit deploy, available
in a couple of minutes

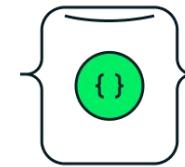


mongoDB® Atlas





MongoDB® Atlas



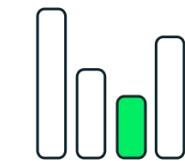
Self-service
and elastic



Global and
highly available



Secure by
default



Comprehensive
monitoring



Managed
backup



Cloud
agnostic

Application Data Platform

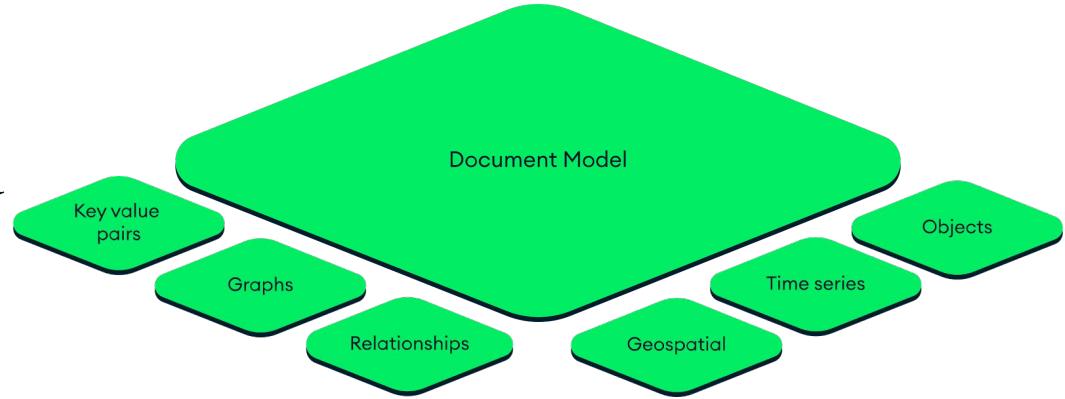


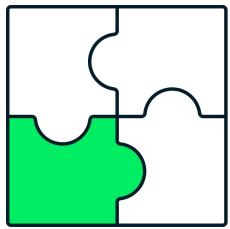


Built around a modern, distributed databases built for developers

Data model that maps to how developers think/code; flexible while allowing data governance when needed

Strongly consistent, support for ACID transactions



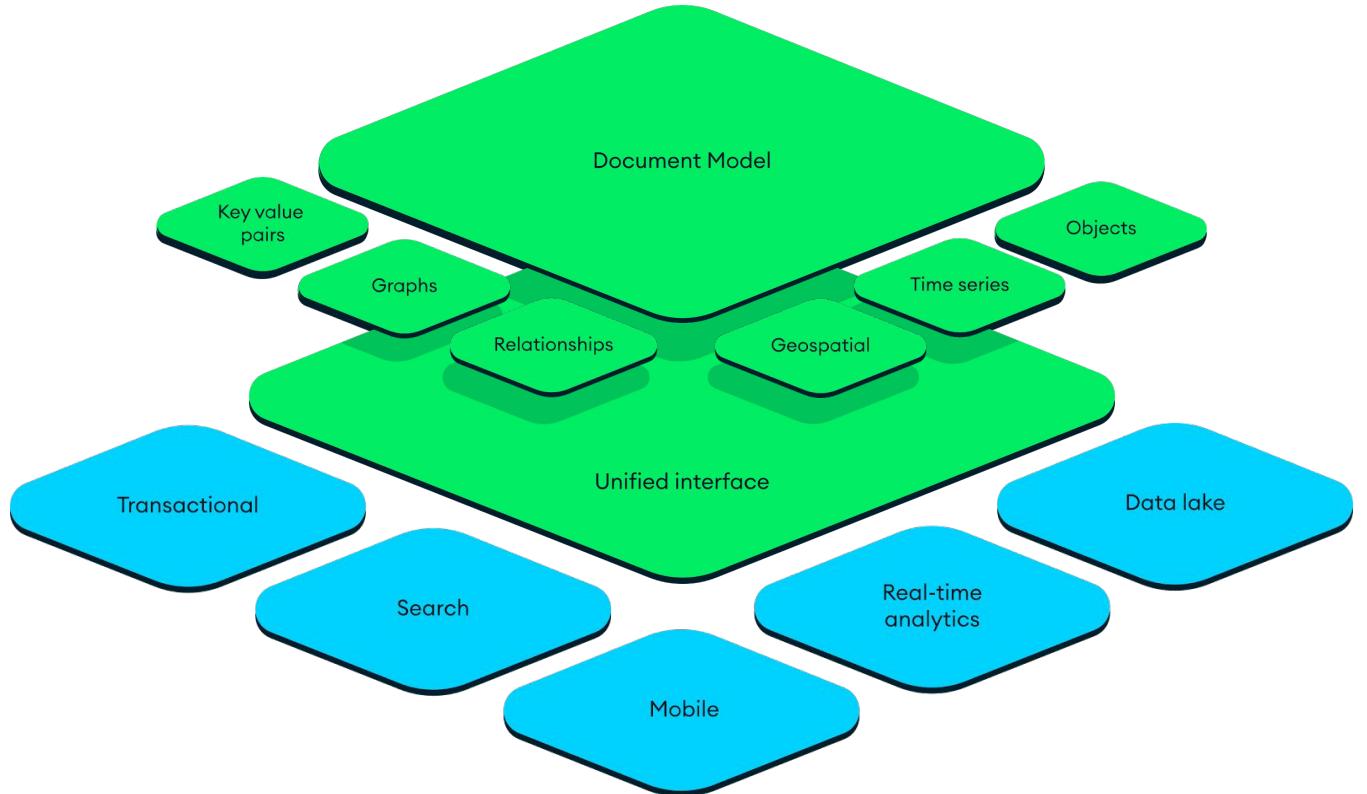


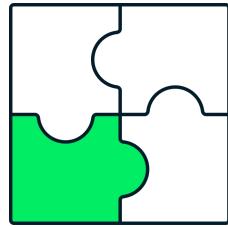
Address a range of application use cases w/o adding complexity

Able to support full-text search functionality for delivering a great user experience

Able to support data on mobile devices at the edge w/o having to manually sync data

Able to deliver real-time analytics on live data w/o having to move data back & forth





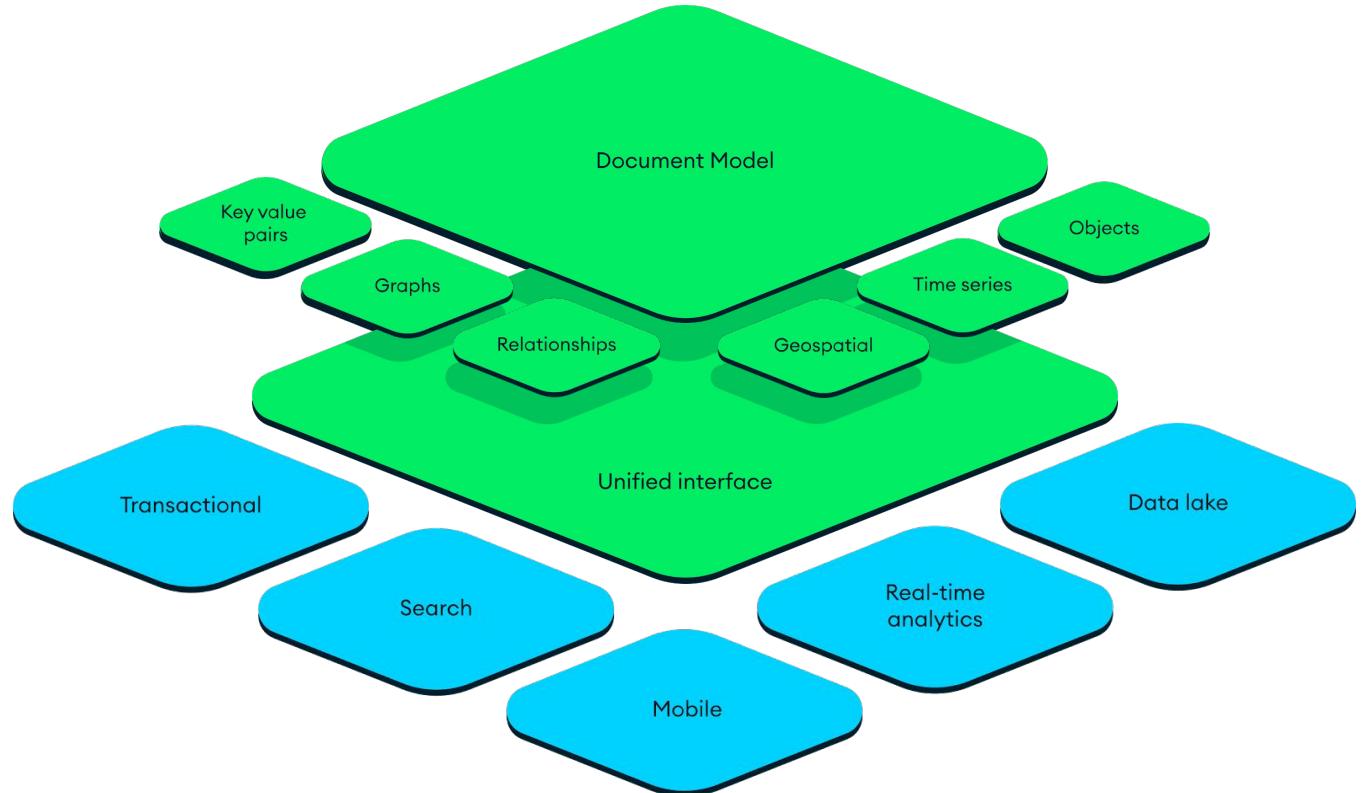
Address a range of application use cases w/o adding complexity

Able to support full-text search functionality for delivering a great user experience

Able to support data on mobile devices at the edge w/o having to manually sync data

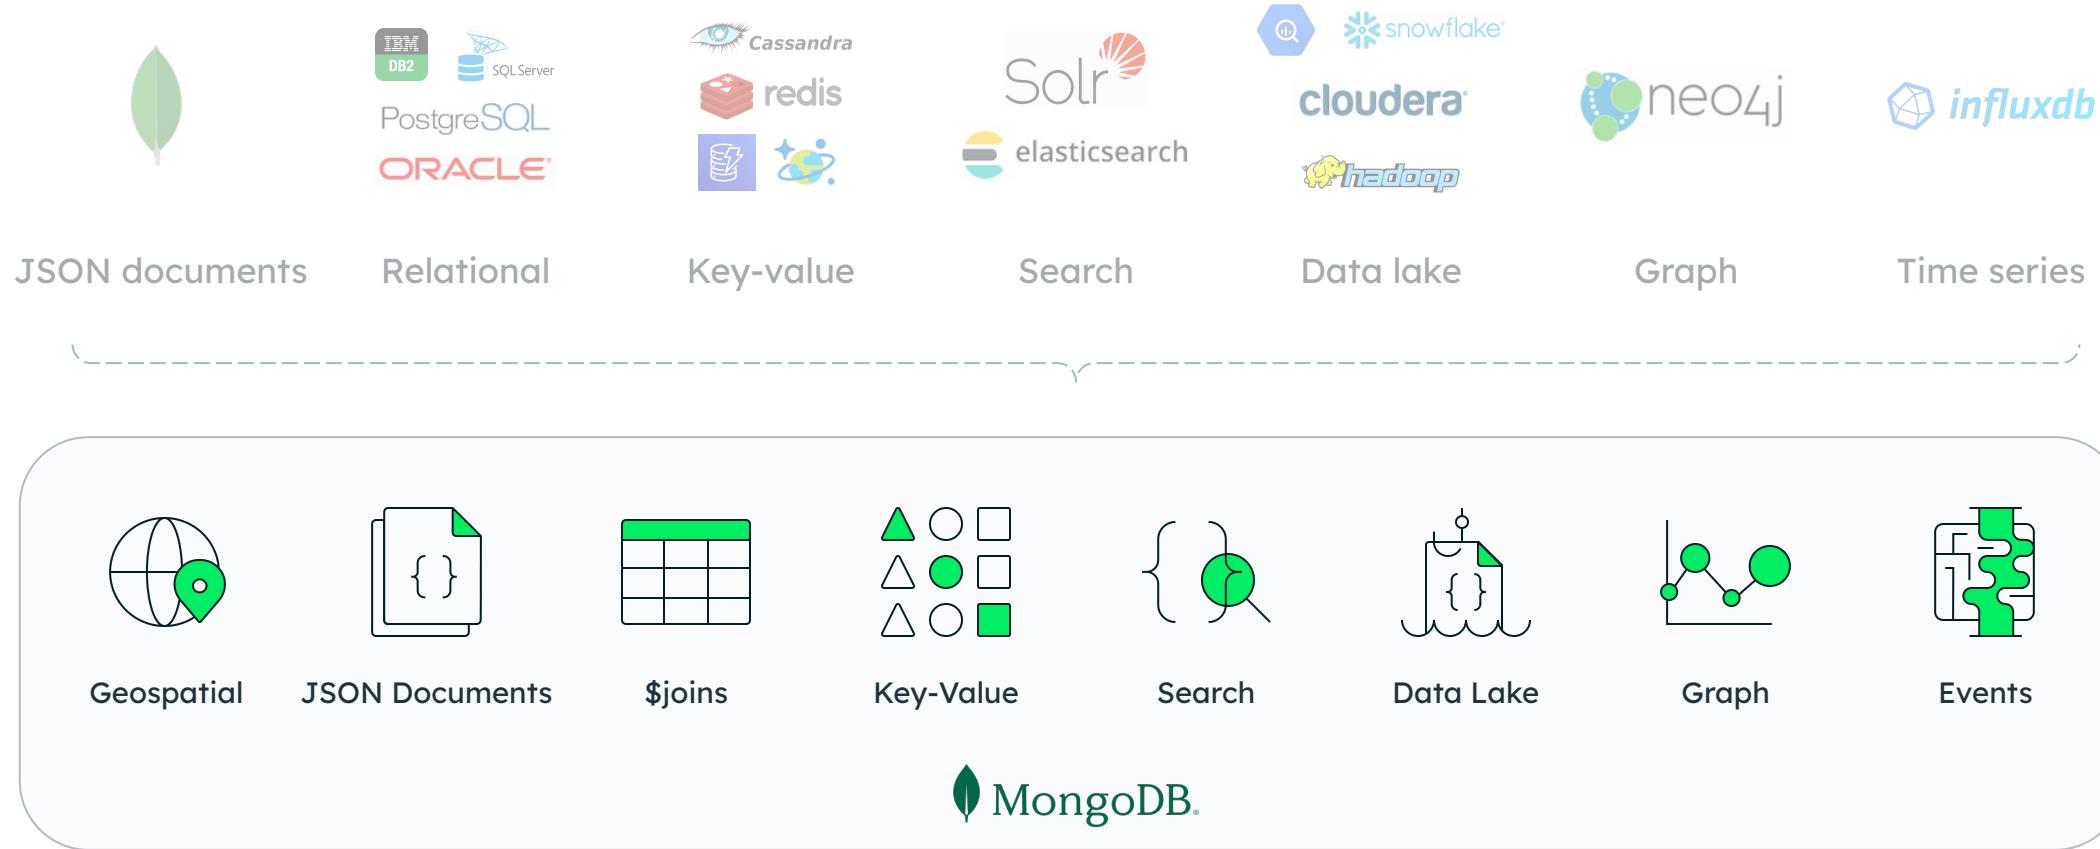
Able to deliver real-time analytics on live data w/o having to move data back & forth

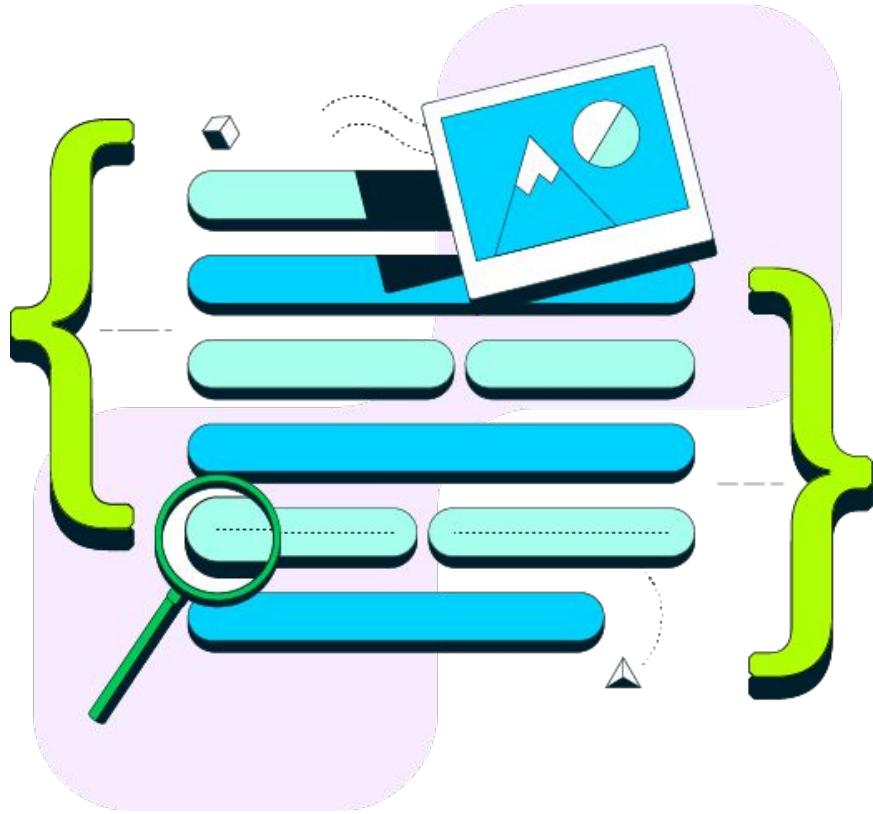
One easy query interface for a variety of workloads





Reduce complexity





Atlas Search

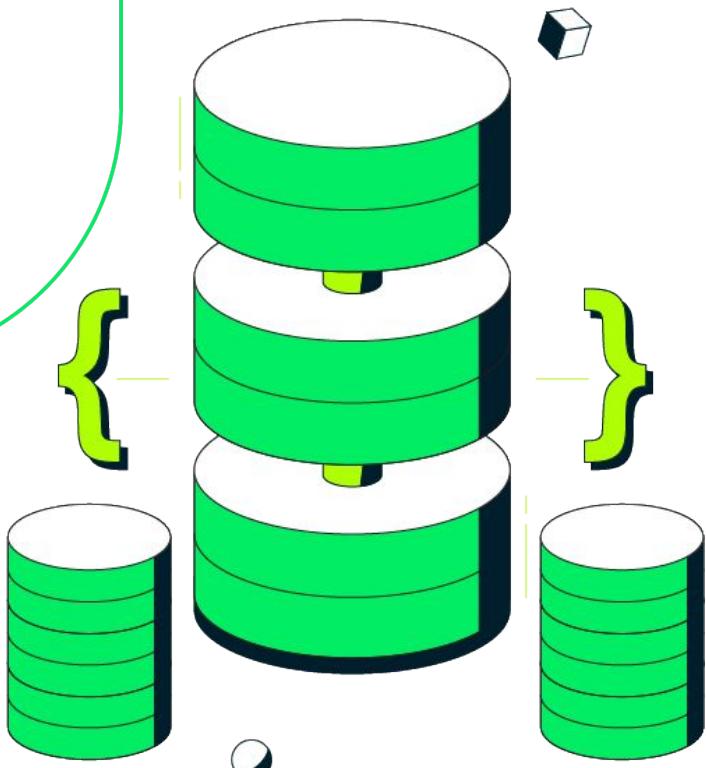
Embedded full-text search in MongoDB Atlas that gives developers a seamless and scalable experience for building relevance-based app features

- Leverage the power of **Apache Lucene**, the world's leading search engine technology
- Eliminate the need to run a separate search system alongside your database

Customers are delivering search features **30%-50% FASTER**



Atlas Data Federation



Analyze cloud data

On-demand query service using the MongoDB API

Perform federated queries

Combine real-time app data with cloud data using a single connection string

Seamlessly work with data

Bring your own AWS S3 buckets so data stays in-place and in its native format



Atlas Charts



Fast to visualize

Built for the document model

No ETL, data movement or duplication required

Easy to start and share

Integrated with MongoDB Atlas

Interactive dashboards with secure sharing

Powerful insights + experiences

No-code aggregation, runs on secondary nodes

Embed charts in apps via IFrame or SDK